

Geometric models reveal that delayed memories are simplified versions of our own earlier recountings

Paxton C. Fitzpatrick¹, Alishba Tahir^{1,2}, Jennifer Xu¹, and Jeremy R. Manning^{1,*}

¹Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755, USA

²ALISHBA AFFILIATION PLACEHOLDER
ALISHBA AFFILIATION PLACEHOLDER

*Corresponding author: Jeremy.R.Manning@Dartmouth.edu

Abstract

What is a second telling of an experience a telling of? Memories change with time and with retrieval, so each recounting may be shaped less by the experience itself than by how we last recounted it. Here we use a geometric framework to compare two recountings of the same experience. We model an experience and each recounting of it as trajectories through a semantic embedding space, then segment each trajectory into events. We applied our approach to data collected as participants watched a television episode, verbally recounted it immediately, and recounted it again one week later. Participants' delayed recountings resembled their own immediate recountings more closely than they resembled the episode itself. Their structure was specific enough to identify the immediate recounting that preceded them. Delayed recountings were also well described as geometrically simplified versions of immediate recountings, in that the events whose removal least distorts a trajectory's shape were those most likely to be dropped. What we retrieve after a delay therefore appears to be a person-specific distillation, rather than the experience as it was first encountered.

18 Introduction

19 Suppose that you attend a friend’s wedding. That evening you describe it to your partner, re-
20 counting how the ceremony ran long, how the best man’s toast was unexpectedly moving, and
21 how a small child wandered onto the dance floor midway through the first song. A week later a
22 coworker asks how the wedding was, and you tell the story again. What is that second telling a
23 telling of? One appealing answer is that both tellings draw on the same underlying memory of
24 the wedding, and that whatever differences arise between them reflect the ordinary imprecision of
25 putting an experience into words. A second possibility, which motivates the work we report here,
26 is that your first telling has itself become part of what you are remembering.

27 What happens to an experience between living it and describing it later? Since our brains
28 do not maintain an exhaustive account of everything we encounter, the very act of perceiving an
29 experience already necessarily entails compressing and distorting it. We attend closely to some
30 aspects of a scene while ignoring others, and we tend to interpret ambiguous moments in light
31 of what has happened previously and we expect will happen next. But once an experience has
32 ended, neither the experience itself nor our moment-to-moment perception of it remains available,
33 to us or to anyone else. What remains is instead the product of a further distillation, in which
34 our perception of the experience is reduced to something our memory systems can store and
35 later reconstruct. Which aspects of an experience survive this second distillation may depend on
36 what we attended to, what we anticipated would matter later, what we found interesting, and the
37 thoughts, motivations, and biases we brought with us. Because remembering entails reconstructing
38 an experience rather than replaying it, the resulting account often preserves the general “shape”
39 of what happened while discarding many of its particulars (Bartlett, 1932; Schacter, 2012; Schacter
40 et al., 2011). The units over which these distillations operate appear to be events, in that people
41 segment continuous experience into discrete episodes at moments of change, and those boundaries
42 shape what is remembered afterwards (Baldassano et al., 2017; Radvansky and Zacks, 2017). In
43 prior work we characterized this transformation geometrically by modeling an experience and a
44 person’s recounting of it as trajectories through a semantic space, and asking which properties of

45 the experience’s trajectory are preserved in the recounting (Heusser et al., 2021).

46 Yet memories do not remain fixed once they have been encoded, pointing to further transfor-
47 mations that a single recounting cannot reveal. Information may be lost over time; experiences we
48 have after an event may color how we later remember it; similar experiences may interfere with
49 one another; and we may simply re-evaluate what happened in light of what we have learned
50 since. The act of retrieving a memory itself appears to be one of the forces that changes it, in that
51 a memory can become labile when it is reactivated and must be stabilized again afterwards (Hup-
52 bach et al., 2007, 2009; Nader et al., 2000; Sederberg et al., 2011). Retrieval also changes what will
53 be accessible later—often substantially—relative to memories that go unretrieved (Karpicke and
54 Roediger, 2007; Roediger and Karpicke, 2006). Further, the particular way in which we recount
55 an experience appears to matter as well, since people who are asked to describe an event for one
56 purpose later remember it in ways that are biased toward that description (Tversky and Marsh,
57 2000). Memories also become more semantic and less perceptually detailed, both with the passage
58 of time and with repeated remembering (Lifanov et al., 2021). Repeated retrieval of naturalistic
59 events over the course of a week has been found to stabilize each person’s recall (for example,
60 which details they include) without causing different people’s recalls to converge (Melega et al.,
61 2025). Taken together, these observations suggest that each time we remember an experience, the
62 substance of the underlying memory we are retrieving may differ from what we retrieved before.

63 How might we compare two recountings of the same experience when the experience and
64 both recountings take the form of unconstrained natural language? Recent work has begun to use
65 language models to quantify the content of what individuals say when they recount what they
66 remember (Fenerci et al., 2025; Fitzpatrick et al., 2026; Heusser et al., 2021; Manning et al., 2022).
67 The approach developed in our prior work embeds the description of an experience, and each
68 individual’s recounting of it, into a common high-dimensional space whose coordinates reflect
69 the concepts under consideration, and through which the experience and recountings each unfold
70 over time as a geometric trajectory (Fitzpatrick et al., 2026; Heusser et al., 2021; Vera et al., 2025).
71 We then segment the experience and each recounting into events and identify which event from
72 the experience each recall event describes, yielding sequences of events that may be compared

73 directly across people and across delays (Baldassano et al., 2017; Chen et al., 2017; Heusser et al.,
74 2021).

75 Here we apply our approach to data from an experiment in which participants viewed an
76 episode of a television show, recounted its events immediately afterward, and then recounted
77 them again following a week-long delay. Because we collected two recountings of the same
78 experience from each participant, we are able to ask what participants' delayed recountings most
79 closely resemble. If retrieval recovered the original experience in an uncorrupted form, then both
80 recountings should be about equally similar to the experience itself. We find instead that the
81 structure of a participant's delayed recounting is more similar to that of their own immediate
82 recounting than to that of the original episode. It is also more similar to their own immediate
83 recounting than to other participants' immediate recountings, to the degree that an individual's
84 immediate recall can be identified from their delayed recall on the basis of this structure alone.
85 What governs retrieval after a delay therefore appears to be a person-specific distillation, rather
86 than the episode as it was first experienced.

87 Quantifying the semantic geometry of individuals' recalls at different points in time offers the
88 opportunity to characterize the transformations undergone by our memories beyond an initial
89 recounting. For example, if individuals' recountings tend to retain their idiosyncratic structures
90 after a delay, how *do* the structures of our memories tend to change with time? There is a particular
91 qualitative change our memories undergo as we naturally forget aspects of experience, implied by
92 common descriptors such as "fading", "blurring", becoming "fuzzy", or "decaying" (Thorndike,
93 1913). But the nature of this change can be difficult to characterize more concretely. Here we show
94 that this transformation is well described as geometrically simplifying the path through semantic
95 space an individual follows as they recount an experience, in that applying existing algorithms
96 for removing detail from a geometric curve to participants' immediate recalls identifies the details
97 they are likely to forget following the delay.

Results

[OUTLINE — placeholder structure, not final prose. Bracketed italic text marks content still to be written. No statistic below has been computed or verified; every “report [STAT]” is a hole to be filled from the cited notebook’s output.]

[FIGURE BUDGET: 8 main-text figures maximum. The plan below uses 6, leaving 2 in reserve (one likely wanted for the RDP analysis if it needs a second panel set, one for contingency). Figure 1 (**paradigm.pdf**) is the only finished main-text figure; the other PDFs already in **paper/figures/source/** are analysis outputs that may feed panels rather than figures in their own right. Nothing else is made yet.]

From episode and recalls to comparable event sequences

[Walk the reader through the pipeline, briefly and conceptually (details live in the Methods), up to the point where we have three things in hand: (a) the sequence of events comprising the episode, (b) each participant’s sequences of immediate and delayed recall events, and (c) the mapping between each recall event and the episode event it describes. This is the setup that everything below depends on — the point to land is that the episode and both recalls now live in a common space and are segmented into comparable units.]

[Figure 1 — **paradigm**. **paradigm.pdf**, already in **paper/figures/source/**. Session-1 row: instructions, episode (24 min), immediate recall (10–60 min), prediction (5 min); one-week gap; session-2 row: instructions, delayed recall (10–60 min), session-2 episode (20/22 min), recall of that episode (10–60 min). Planned revision: reduce the opacity of the blocks for the phases we don’t analyze — the session-1 prediction block, and the session-2 episode-viewing and subsequent-recall blocks — so the reader’s eye lands on the immediate/delayed recall pair that the paper is actually about.]

[Figure 2 — **methods overview**. To be made. Annotations and recall text → trajectories → event-segmentation procedure → event sequences. Loosely modeled on Fig. 2 of Heusser et al. (2021), extended with the (new) event-segmentation step. Layout still open. Candidate source material already generated: **episode-corrmat.pdf** and **k-optimization-episode.pdf** (from **event-segmentation.ipynb**), and the per-recall similarity heatmaps and match matrices **atlep1-and-delayed-recall-similarity-heatmaps.pdf**

124 / -matchmats.pdf.]

125 **Participants recall the episode's events in a similar order before and after the** 126 **delay**

127 *[Motivation: we wanted to know whether participants remembered the episode similarly before and after the*
128 *delay. The simplest way to ask this is whether they recalled the same events in the same order.]*

129 *[Result: recall-order similarity. Computed in `self-similarity-analyses.ipynb`, section “Is de-*
130 *layed recall order more similar to own immediate recall order than to others’ immediate recall orders?”. The*
131 *measure compares pairwise precedence relations over the episode events shared between two recalls, then*
132 *percentile-ranks each participant’s delayed-vs-own-immediate similarity against their delayed-vs-others’-*
133 *immediate similarities, tested against chance (0.5) with a Wilcoxon signed-rank test. Report [STAT].]*

134 *[Figure 3 — recall-order similarity. To be made; multi-panel.]*

135 *[Panel A: illustration of the method. For one example participant, show the matrix of pairwise recall-*
136 *order-precedence values over pairs of episode events — probably just the lower triangle — for their delayed*
137 *recall beside their own immediate recall. Grey out (or reduce the opacity of) cells for event pairs not recalled*
138 *in both sessions, since the comparison is restricted to pairs present in both. Then show the same thing again*
139 *for that participant’s delayed recall against a different participant’s immediate recall, which is the contrast*
140 *the analysis actually tests. This maps directly onto `pairwise_precedence()` and the `shared_events`*
141 *argument to `order_similarity()` in the notebook, so the panel and the code should be checked against*
142 *each other.]*

143 *[Panels B–D: de-sloped versions of panels B–D of the reference mock-up. Include the analogues of that*
144 *mock-up’s B and D so that the same panel types can reappear in the recall-structure figure (Figure 5), letting*
145 *the reader compare order and structure on identical terms — this parallelism is the reason for the pairing*
146 *and should be preserved in the final layout.]*

147 *[Note on ordering: this is the least novel analysis of the three and the one least dependent on our*
148 *framework — order could be measured without any embedding model. It goes first precisely because it*
149 *establishes the basic effect on the simplest possible terms, setting up the contrast with what follows.]*

150 **The semantic structure of participants' recalls is preserved across the delay**

151 *[Motivation: order is only the coarsest description of a recall. Because our approach embeds recalls in a*
152 *continuous semantic space, we can ask a stronger question — whether the structure of how a participant*
153 *moved through the episode's content is preserved, not merely the sequence in which they hit the events.]*

154 **[Figure 4 (optional) — 2D trajectories.** *Decision deferred, and it depends on how this section's*
155 *narrative settles. Argument for including it: it introduces the idea of geometric comparison and establishes*
156 *the 2D paths before they get used as examples in Figure 5, so the reader isn't meeting the representation*
157 *and the result at the same time. Argument against: it spends one of 8 figure slots on setup. Middle*
158 *option — a compact version showing only the episode path plus immediate and delayed paths for a handful*
159 *of participants, with the full version relegated to the supplement. Source: 2d-trajectories.ipynb*
160 *(atlep1-paths2d.pdf and delayed-paths2d.pdf, imm-del-2d.pdf); note those savefig calls are*
161 *currently commented out in the notebook.]*

162 *[Result: trajectory/event-structure similarity. Computed in self-similarity-analyses.ipynb,*
163 *sections "Is delayed recall more similar to own immediate recall than to episode?" and "Is delayed recall*
164 *more similar to own immediate recall than to others' immediate recalls?". Both use DTW warp cost (average*
165 *cosine distance between matched events) over mean-centered event sequences; the second percentile-ranks*
166 *self against others and tests against chance. Report [STAT] for each.]*

167 *[Framing: this result appears strong enough to present as a "fingerprint" — i.e., that one can identify*
168 *which immediate recall belongs to a given participant from the structure of their delayed recall alone. Frame*
169 *it that way in both the text and the figure, but confirm the effect size supports the claim before committing*
170 *to that language.]*

171 **[Figure 5 — recall-structure similarity.** *To be made; the largest figure in the paper, covering both*
172 *analysis (A) "delayed is more similar to immediate than to the episode" and analysis (B) "delayed is more*
173 *similar to own immediate than to others' immediates". Layout still open.]*

174 *[Example trajectory panels (3, from the current draft): the left panel is common to both analyses, so*
175 *it should sit centrally in the final arrangement; the right panel is the contrast for analysis (A) and the*
176 *middle panel is the contrast for analysis (B). Planned refinements: thicker lines; faint connectors between*

177 DTW-matched events within each panel (already present in the draft); possibly PCHIP splines rather than
178 straight segments between points, to make the paths read as curves. Under consideration — a fourth panel
179 showing pairwise mean point-by-point distances between trajectories, which would require upsampling each
180 trajectory to a common number of points and comparing them in order (or substituting DTW cost, if that
181 turns out to read better).]

182 [Analysis (A) panel: per-participant connected dots showing each participant’s delayed-to-immediate
183 versus delayed-to-episode alignment cost. Probably without the KDEs from the draft version.]

184 [Analysis (B) panels (3): de-slopified versions of the current draft. Two of them are the direct analogues
185 of Figure 3’s panels B and D and show stronger effects — keeping the panel types visually identical across
186 the two figures is what makes “structure beats order” legible without the reader doing arithmetic. The third
187 gets at how person-specific the similarity is, and is the panel that carries the “fingerprint” claim.]

188 [Candidate supporting material: `atlep1-recall-cormats.pdf` and `delayed-recall-cormats.pdf`
189 (from `event-segmentation.ipynb`).]

190 **Delayed recalls are simplified versions of participants’ immediate recalls**

191 [Motivation: having established that a participant’s delayed recall structure resembles their own immediate
192 recall structure, the natural next question is how it changes over the delay.]

193 [Result: trajectory simplification. Computed in `rdp-analyses.ipynb`. Each participant’s immediate
194 recall trajectory is simplified with the Ramer-Douglas-Peucker algorithm over a range of tolerances (ϵ), and
195 the ϵ minimizing DTW warp cost to their delayed recall trajectory is selected. Sub-results to report: (i) the
196 relationship between how many events a participant dropped across the delay and how many RDP removes
197 at the best-fitting ϵ (Kendall’s τ); (ii) the same against ϵ itself; (iii) whether RDP predicts which events
198 were forgotten, benchmarked against a hypergeometric null and compared to predictions derived from other
199 participants’ trajectories. Report [STAT] for each.]

200 [Figure 6 — RDP simplification. To be made, and how best to illustrate the simplification is genuinely
201 unresolved — flagged as an open design question rather than a layout still to be drawn. The regression plots
202 for sub-results (i) and (ii) exist in `rdp-analyses.ipynb` but have no `savefig` calls, so nothing is exported
203 to `paper/figures/source` yet. One thing worth trying: since Figure 5 will already have taught the reader

204 *to read 2D trajectories, showing one participant’s immediate path beside its RDP-simplified version and*
205 *their actual delayed path would make “the delayed recall looks like a simplified version of the immediate one”*
206 *visible directly, with the quantitative panels as support. If the analysis needs a second panel set, one of the 2*
207 *reserve figure slots can absorb it.]*

208 *[@paxton: does the simplification story want a fourth subsection on what the dropped events have in*
209 *common (e.g. are they lower-salience, more peripheral to the narrative), or is that a Discussion point?]*

210 Discussion

211 We asked participants to watch an episode of a television show, to recount what they remembered
212 immediately afterwards, and to recount the episode again following a week-long delay. Embedding
213 the episode’s annotations and each recounting into a common semantic space allowed us to
214 represent all three as trajectories, to segment each trajectory into events, and to compare them on
215 common terms. Participants recalled the episode’s events in a similar order before and after the
216 delay. The semantic structure of a participant’s delayed recall was more similar to the structure
217 of their own immediate recall than to that of the episode, and it was more similar to their own
218 immediate recall than to other participants’ immediate recalls. Delayed recalls were also well
219 described as simplified versions of immediate recalls, in that removing detail from a participant’s
220 immediate recall trajectory brought it closer to the delayed recall trajectory that followed it.

221 Taken together, these results suggest a way of thinking about what recounting an experience
222 does to it. If experiences and memories may be described as trajectories through a semantic
223 space, then a recounting is a traversal of one of those trajectories (Heusser et al., 2021). Semantic
224 space differs from physical space in that nothing prevents a narrator from jumping between distant
225 locations, and yet participants’ recountings turn out to be strongly autocorrelated in practice, which
226 is the regularity that our event-segmentation approach relies on. Because movement through
227 the space is relatively continuous, describing how one arrives at a later event from an earlier
228 one implies that the intervening region was traversed in some sense, whether it was articulated
229 explicitly or left for the listener to fill in. Some events therefore behave like control points, in that

230 removing one substantively changes the shape of the path. Others behave like waypoints, which
231 fall close to the line between their neighbors and may be removed with little consequence, since
232 travelling between the surrounding control points already implies passing near them. Trajectory
233 simplification provides a way of telling the two apart, and this distinction appears to hold at both
234 of the transitions we can examine. The events that simplification removes from the episode's own
235 trajectory are the events that the fewest participants recounted immediately afterwards ([**VERIFY**
236 **THIS**]), and the events that simplification removes from a participant's immediate recall trajectory
237 are the ones that participant tended to drop across the delay.

238 That not all events are equally memorable is well established. Gist and central details, for
239 example, survive delays that perceptual and peripheral details do not (Gilboa and Moscovitch,
240 2021; Heusser et al., 2021; Sekeres et al., 2016, 2018), and autobiographical narratives shed some
241 classes of content faster than others (Diamond and Levine, 2020; Su et al., 2026). Memories also
242 become more semantic and less perceptually detailed with the passage of time and with repeated
243 remembering (Lifanov et al., 2021), and events with stronger connections to the events around them
244 are better remembered than more isolated ones (Antony et al., 2024; Lee and Chen, 2022). What our
245 results add is a change in where that asymmetry comes from. Centrality is ordinarily established
246 from the outside, by asking raters which details matter to the plot, or by sorting content into
247 categories fixed in advance. Here it follows from the geometry of an individual's own trajectory,
248 and it is computable from that trajectory alone. An event is central because of where it sits relative
249 to the rest of the path, and not because a rater judged it to be important. This turns the observation
250 that some events are more memorable than others into a candidate explanation of why.

251 Why should waypoints be the events that go? One possibility is that they are redundant,
252 in the sense that a listener, or the rememberer's own future self, can recover them from the
253 events on either side. Another is that they are simply weaker in themselves, for example less
254 distinctive, less emotionally charged, or less causally connected to the rest of the narrative (Antony
255 et al., 2024). These accounts make similar predictions about which events will be forgotten, but
256 they differ in what those predictions should depend on. If events are dropped because they are
257 intrinsically weak, then a simplification derived from any participant's trajectory should predict

258 forgetting about as well as one derived from the participant's own, since weak events would be
259 weak for everyone. If they are dropped because they are recoverable from their neighbors, then
260 a participant's own trajectory should do better, since control-point status is defined relative to the
261 rest of that person's path. Our results favor the latter account. Read this way, dropping a waypoint
262 is an economy rather than a loss, and it is worth noting that events which go unreported are
263 not always unavailable, since a shift in perspective at retrieval can recover details that an earlier
264 recounting left out (Anderson and Pichert, 1978; Oedekoven et al., 2017).

265 The same argument implies that what counts as a control point should be idiosyncratic. An
266 event's status depends on the path it sits within, and different people who watched the same episode
267 construct different paths through it, so the events whose removal would deform one person's
268 trajectory need not be the events that matter for another's. This is what we take the specificity of
269 the delayed recalls to mean. That a delayed recall can be matched to the immediate recall which
270 preceded it, on the basis of structure alone, is a fingerprint in the sense that the term is used for
271 the neural signatures by which individuals may be identified from functional connectivity or from
272 their responses to naturalistic stimuli (Anderson et al., 2020; Finn et al., 2020, 2015). The parallel
273 lies in the identification logic rather than in the measurement, since ours is a behavioral signature,
274 recovered from what a participant said rather than from how their brain responded. People differ,
275 for example, in where they place event boundaries in a continuous experience (Sava-Segal et al.,
276 2023), in how much of their recall structure they share with others (Chen et al., 2017; Nguyen
277 et al., 2019), and in the character of their autobiographical remembering more generally (Palombo
278 et al., 2018). Repeated retrieval of naturalistic events has been found to stabilize each person's
279 recall without causing different people's recalls to converge (Melega et al., 2025), which is what
280 one would expect if each person is stabilizing their own set of control points.

281 Everything we measured came from a verbal recounting, and so a reasonable objection is that
282 control points and waypoints are properties of storytelling rather than of memory. We think that
283 dichotomy is less clean than it first appears. Recounting is the only observable available in a study
284 of this kind, so the two are confounded by construction, and no analysis of transcripts alone can
285 fully separate them. More to the point, the way a person tells a story is not merely a window onto

286 what they remember, but also an input to what they will remember next. Retelling an event for a
287 particular audience or purpose, for example, biases how it is recalled later (Marsh, 2007; Tversky
288 and Marsh, 2000), and reproducing a narrative repeatedly reshapes it in ways that persist (Bergman
289 and Roediger, 1999; St Jacques and Schacter, 2013). If the function of recounting is to let a listener,
290 or one's own future self, re-traverse the trajectory of the original experience, then communicative
291 efficiency and mnemonic efficiency are not competing explanations so much as the same pressure
292 seen from two sides.

293 We acknowledge several limitations in the current study. First, every participant recalled the
294 episode immediately, and so we cannot separate two accounts of what we observed. A delayed
295 recall may resemble the immediate recall because the immediate recall was itself re-encoded and
296 then retrieved a week later, or because a single stable memory was sampled twice and each
297 sampling was shaped by the same person-specific distillation. Our measures are symmetric and
298 cannot establish direction, and the design lacks an arm in which the immediate recall is withheld.
299 Designs that do include such a control, in which retrieved information is later found nearer
300 to what was retrieved than to what was originally presented, offer the strongest evidence that
301 retrieval writes its own output back into the trace (Bridge and Paller, 2012; Nader and Hardt, 2009;
302 Nader et al., 2000). Second, we used one episode, one delay, and one intervening recounting, and
303 it remains to be seen how the picture changes across stimuli, across longer intervals, or across
304 many successive retellings (Chan et al., 2009; Lifanov et al., 2021). Third, control-point status is
305 defined with respect to one particular simplification algorithm, and a different algorithm might
306 well partition the same events differently. Fourth, our procedure matches every recall event to
307 an event in the episode, so content that a participant introduced but that does not appear in the
308 episode falls outside what we are able to measure. Fifth, our measure of what connects one event
309 to another is built from semantic similarity, and so it cannot on its own distinguish events that are
310 close in meaning from events that are causally linked (Antony et al., 2024; Lee and Chen, 2022; Xu
311 et al., 2024).

312 Speculatively, one way to think about how this might work is in terms of attractor dynam-
313 ics (Hopfield, 1982). Rather than supposing that we hold literal duplicate copies of a memory,

we might imagine that laying down a memory carves out a basin in an energy landscape, and that remembering the experience deepens a basin near wherever we happen to land. Combined with the observation that remembering distorts, this implies that what gets re-encoded on each occasion may not be the original memory, but a distorted version of it. Successive retrievals might then arrive at a basin corresponding to the original experience, or at one corresponding to any of our previous rememberings of it. A related point has been made from the perspective of context reinstatement, namely that each time we remember a prior experience, our memory of that experience becomes colored by the contextual thoughts we were having at the time of each remembering (Manning, 2024). We did not test this picture, and our data cannot distinguish it from other mechanisms that would leave the same trace in a transcript, but it offers a natural reading of why a delayed recounting should resemble an earlier recounting more closely than the experience that both describe.

Whatever the underlying mechanism, our account suggests what might become of a memory that is recounted many times. If each retelling preserves control points and sheds waypoints, and if each retelling then becomes the trajectory that the next one simplifies, a memory might converge over many retellings toward a skeleton of control points, with the intervening detail increasingly implied rather than stated. Whether that convergence has a floor, and whether the surviving skeleton continues to resemble the original experience or drifts steadily away from it, are questions our data cannot answer. The possibility seems worth taking seriously outside the laboratory, since the settings in which people are asked to recount an experience repeatedly are often the settings in which accuracy matters most. An eyewitness who has already described what they saw, for example, is describing a description the next time they are asked (Chan et al., 2009). A patient recounting a formative episode across successive sessions may be reshaping it in the process, and a teacher who explains an idea the same way each term may be simplifying it past the point where it still carries what made it worth explaining.

Ultimately, our work points toward a set of questions about how the geometry of a recounting relates to the experience that it describes. Do the events that behave like control points in a recounting correspond to the boundaries at which neural activity shifts during the experience

342 itself (Baldassano et al., 2017; Chen et al., 2017)? Can they be predicted in advance from the
343 stimulus, or are they recoverable only after the fact, from a particular person's path? Does a listener
344 who hears a recounting actually re-traverse the speaker's trajectory, filling in the waypoints that the
345 speaker left out? Would causal and importance judgments, collected for this same episode, show
346 that geometric control points and causally pivotal events are one notion arrived at by different
347 routes? Questions of this kind become tractable once what a person retains from an experience
348 is treated as a shape that can be measured, rather than as a quantity that can only be counted.
349 Because that shape may be recovered from what someone says, and because control points are
350 identifiable from a single recounting, it may eventually be possible to ask what a student will still
351 carry from a lecture, or what a patient will still carry from a conversation with their physician,
352 before the delay that would reveal it has passed. We hope that this offers a way of thinking about
353 forgetting as something other than loss, in that what falls away between one telling and the next is
354 largely what the remaining structure already implies, and what survives is the shape in which an
355 experience can be carried forward and handed on.

356 **Methods**

357 Our experimental protocol was approved by the Committee for the Protection of Human Subjects
358 at Dartmouth College. We obtained written informed consent from all participants prior to the
359 start of both experiment sessions.

360 **Participants**

361 We enrolled a total of 70 Dartmouth undergraduate students in our study. Participants received
362 optional course credit for enrolling and were pre-screened to ensure they had not previously seen
363 any episode of either TV show viewed during the experiment. No statistical method was used
364 to predetermine sample size. Prior to data processing and analysis, we excluded a total of 17
365 participants from our dataset: 4 encountered technical issues with the experiment, 2 self-reported
366 difficulty with the task due to illness, 4 spoke for less than 10 minutes during the immediate recall

367 phase, and 7 did not return for the second experiment session. The data we analyzed comprised
368 the remaining 53 participants.

369 At the start of the experiment, we asked each participant to complete a demographic survey
370 that included questions about their age, sex, race, and ethnicity. Participants' ages ranged from 18
371 to 22 years (mean: 18.91 years; standard deviation: 1.01 years). A total of 41 participants reported
372 their sex as female and 12 participants reported their sex as male. Participants reported their races
373 as White (34 participants), Asian (16 participants), Black or African American (3 participants),
374 American Indian or Alaska Native (1 participant), and other (1 participant). (Note that some
375 participants selected multiple racial categories.) A total of 52 participants reported their ethnicity
376 as "Not Hispanic or Latino" and 1 participant reported their ethnicity as "Hispanic or Latino."

377 We also asked participants their native spoken language; whether they had any hearing, speech,
378 or color vision impairments; and about any current medications or recent injuries that could impact
379 attention or memory. All participants reported that their native spoken language was English, that
380 they had no hearing or speech impairments, and that they had normal color vision. A total of 52
381 participants reported no attention- or memory-relevant medications or injuries. One participant
382 reported sustaining a concussion more than 3 years prior to the experiment and was not excluded
383 from our dataset.

384 Additionally, before each of the two experiment sessions, we asked participants about their
385 sleep, coffee consumption, and level of alertness. Participants reported having slept between 5
386 and 11 hours the night prior to session 1 (mean: 7.37 hours; standard deviation: 1.45 hours) and
387 between 5 and 11 hours the night prior to session 2 (mean: 7.39 hours; standard deviation: 1.23
388 hours). They reported having consumed between 0 and 3 cups of coffee on the day of session
389 1 (mean: 0.43 cups; standard deviation: 0.72 cups) and between 0 and 3 cups of coffee on the
390 day of session 2 (mean: 0.48 cups; standard deviation: 0.68 cups). Participants reported their
391 current levels of alertness on a 5-point Likert scale. At the start of session 1, 1 participant reported
392 feeling "Very sluggish," 14 reported feeling "A little sluggish," 18 reported feeling "Neutral," 15
393 reported feeling "A little alert," and 5 reported feeling "Very alert." At the start of session 2, 3
394 participants reported feeling "Very sluggish," 9 reported feeling "A little sluggish," 23 reported

395 feeling “Neutral,” 17 reported feeling “A little alert,” and 3 reported feeling “Very alert.”

396 @jeremy: *do we want to report class years and majors? And/or do we want to report post-experiment*
397 *survey responses (how engaging did you find the episode, how easy/difficult was it to follow, how well do you*
398 *think you recalled it, how well do you think you learned characters’ names)?*

399 Experiment

400 The experiment took place over two sessions lasting approximately one hour each, separated by
401 6–8 days (mean: 6.76 days; standard deviation: 0.73 days). Prior to the start of the experiment,
402 participants were assigned to one of two groups (*A* or *B*), alternating based on enrollment order.
403 During session 1, all participants first viewed the pilot episode of the FX series *Atlanta* (“The Big
404 Bang”; duration: 24 min 15 s). Immediately after viewing, participants were given up to one hour
405 to recount what happened in the episode to the best of their abilities. Following Chen et al. (2017),
406 participants were instructed to describe the episode in as much detail as possible and to try to
407 speak for a minimum of 10 minutes. They were told to try to recount the episode’s events in their
408 original order, but that they were free to go back and describe earlier events should they realize
409 they had missed something. After the immediate recall phase, participants were then given up to
410 5 minutes to predict what they thought would happen in the show’s next episode.

411 Upon returning for session 2, all participants were again given up to one hour to recount the
412 *Atlanta* episode they had viewed the week prior. Following the delayed recall phase, participants
413 in Group *A* then viewed the second episode of *Atlanta* (“Streets on Lock”; duration: 21 min
414 57 s) while participants in Group *B* instead viewed the pilot episode of the FOX series *Arrested*
415 *Development* (“Pilot”; duration: 20 min 37 s). Both groups were then given up to one hour to
416 recount the just-watched episode in the same manner and given the same instructions as in session
417 1. We implemented the experiment using the psiTurk platform (Gureckis et al., 2016) and jsPsych
418 library (de Leeuw, 2015), with a custom jsPsych plugin we developed in prior work for recording
419 verbal recall audio (Ziman et al., 2018). Here we analyze participants’ immediate and delayed
420 recalls of the *Atlanta* episode viewed in session 1.

Stimulus annotation and recall transcription

To capture the evolving semantic content of the *Atlanta* episode, we created detailed text annotations of the on-screen content during each individual camera shot throughout the episode (mean annotation duration: 3.05 s; standard deviation: 1.80 s). Specifically, we used ANVIL (Kipp, 2001) to record, for each camera shot, a set of semantically relevant features. These included a brief description of physical occurrences and characters’ explicit actions, a brief description of characters’ internal states (e.g., emotions, motivations) and other implicit content, the names of all characters appearing on screen, a transcription of any dialogue, the name(s) of the character(s) speaking, the shot’s setting, whether the shot took place indoors or outdoors, and whether or not there was music playing during the shot. This feature set was adapted from a similar annotation corpus created by Chen et al. (2017), which we used in developing our trajectory-analysis framework in prior work (Heusser et al., 2021).

We then transcribed each participant’s immediate and delayed verbal recall of the *Atlanta* episode. Our transcription process consisted of three steps: an initial manual transcription, a two-pass prose-formatting conversion with a large language model, and a final manual cleanup and correction step. First, we manually transcribed each recall audio file collected during the experiment. These initial transcripts reflected participants’ spoken recalls verbatim, without any punctuation or additional formatting, and including disfluencies, stuttered repetitions, and other artifacts common to unrehearsed speech. However, since transformer models (such as we used to obtain text embeddings of participants’ recalls) are trained almost exclusively on written text, structural conventions like sentence breaks, punctuation, and capitalization provide an important signal these models leverage to contextual other tokens (e.g., words) and construct a representation of the text’s overall meaning. Meanwhile, speech-native phenomena such as stuttered repetitions and filler words (“like”, “um”, etc.) that are comparatively rare in written text contribute out-of-distribution noise that can be expected to degrade embedding fidelity, and prior work has shown that automated removal of these phenomena improves performance on downstream tasks that assume clean, written-style input (Chen et al., 2022; Wang et al., 2010). We therefore used a large

language model (Qwen3.5 122b; Qwen Team, 2026) to convert participants’ recall transcripts to a structure and style more typical of written text. We used the `langchain_dartmouth` Python package (Stone et al., 2025) to implement this conversion as two sequential passes over each transcript. In the first pass, we prompted the model to split the transcript into sentences at natural breaking points, add punctuation including quotes around reported speech, remove clear non-word disfluencies (“um,” “uh,” “hmm,” etc.), collapse stuttered verbatim repetitions, and mark self-corrections and restarts with em-dashes. In the second pass, we prompted the model to remove all instances of “like” used as a filler word. We removed filler “like” instances in a separate pass because of their high frequency of occurrence in the transcripts and the additional effort required to distinguish them from meaningful uses of “like.” The full texts of our prompts for both passes can be found in *Supplementary methods*. Finally, we manually reviewed all LLM-formatted transcripts to identify and correct any errors, ensure sentence breaks were appropriate and consistent across participants, confirm that no content was improperly added to or removed from the transcripts, and resolve any ambiguities based on the prosody of participants’ original recall audio.

Analysis

Statistics

Wilcoxon tests were one-sided because our alternative hypothesis was directional (delayed matches own immediate *better than* the episode, others’ immediates, etc.).

Constructing episode and recall trajectories

We adapted a method we developed in prior work (Heusser et al., 2021) to transform the *Atlanta* episode and participant’s recalls of it into trajectories through a common representation space using a text embedding model. Briefly, text embedding models take in natural language text of an arbitrary length and output a fixed-length vector representation that captures the text’s semantic meaning. Intuitively, an output vector of length n can be viewed as a coordinate in an

473 n -dimensional text embedding space. While specific model architectures and training paradigms
474 vary, modern text embedding models are broadly trained to output similar coordinate vectors (in
475 terms of cosine similarity, dot product, inverse Euclidean distance, or other geometric measures) for
476 semantically similar input texts, even when those texts share few or no specific words. We obtained
477 text embeddings for the episode and recalls using EmbeddingGemma (Vera et al., 2025), a pre-
478 trained 308-million-parameter encoder-only transformer model that outputs coordinate vectors in
479 a 768-dimensional embedding space.

480 To construct the trajectory for the *Atlanta* episode, we first converted our tabular annotations
481 of the content present during each camera shot to a plain-text format. Separately for each shot,
482 we concatenated the annotation text describing explicit on-screen occurrences and characters’
483 implicit internal states. We then appended “Speech: ” followed by the shot’s dialogue, “Character
484 speaking: ” followed by the name of the speaker, and “Setting: ” followed by the shot’s setting. If
485 any annotated fields were empty for a given shot (e.g., if no character was speaking at that time), we
486 omitted them and their prefixes from the concatenated text. Next, we used EmbeddingGemma to
487 transform the text snippet for each camera shot into a text embedding vector, resulting in a number-
488 of-shots (478) by number-of-features (768) matrix. Since the camera shots varied in duration, we
489 then resampled this series of embedding vectors to a consistent temporal resolution. We assigned
490 each vector to the timepoint midway between the start and end of its corresponding camera shot
491 (in seconds), and used linear interpolation to estimate a vector for each second of the episode. This
492 yielded a number-of-seconds (1,455) by number-of-features (768) matrix—i.e., a 768-dimensional
493 *trajectory*.

494 We then created similar trajectories for each participant’s immediate and delayed recall of the
495 episode. We began by splitting each recall transcript into a list of sentences, with the constraint
496 that splits could not occur within reported speech (i.e., if a participant recounted a multi-sentence
497 quote by a character in the episode, we treated the full quote as a single sentence-unit). We
498 then parsed these sentences into overlapping sliding windows that spanned up to 5 consecutive
499 sentences each. These windows grew to and shrank from their full size at the beginning and end of
500 a participant’s recall sequence—in other words, the first window contained only the first sentence

of the participant’s recall, the second contained the first two sentences, and so on. This ensured that all sentences from a participant’s recall transcript appeared an equal number of times (5) in their sequence of sliding windows. Finally, we used EmbeddingGemma to transform each sliding window into a text embedding vector. Repeated for all recall transcripts, this yielded a separate trajectory characterizing each participant’s recall from each of the two experiment sessions. The number of samples in participants’ recall trajectories (i.e., the number of sliding windows parsed from their recall transcripts) ranged from 61 to 308 (mean: 136.68; standard deviation: 49.70).

Relating episode and recall embedding vectors

In prior work using our trajectory-analysis framework (e.g., Fitzpatrick et al., 2026; Heusser et al., 2021; Manning, 2021; Manning et al., 2022; Raccah et al., 2024), embeddings of the stimulus and participants’ responses were constructed using a topic model (Latent Dirichlet Allocation; Blei et al., 2003) trained solely on the stimulus content. Here we elected to instead use a pre-trained transformer-based text embedding model due to the richer and more nuanced textual information these models are able to capture. (For example, these context-sensitive models can disambiguate the texts “Earn looked at Alfred” and “Alfred looked at Earn” in their representations, whereas bag-of-words models like LDA would represent the two identically.) However, these richer representations also introduce certain geometric nuisances we needed to account for in order to effectively relate the contents of participants’ recalls to that of the episode they viewed. First, because modern transformer models are pre-trained on large and diverse text corpora, the representational capacities of their embedding spaces span an enormous breadth of semantic content. Consequently, in preliminary analyses we found that the embeddings of our annotations and participants’ recalls (which were all “about” the events, themes, and characters of a single *Atlanta* episode) were strongly anisotropic—i.e., they all occupied the same relatively narrow angular region of the model’s overall embedding space. Additionally, because transformer models are sensitive to the syntactic structure of text, their embedding representations carry latent information about the text’s linguistic register in addition to its semantic content. Accordingly, we found there was a consistent angular offset between embeddings of the episode annotations and those of par-

528 ticipants’ recalls, reflecting the fact that the annotations comprise grammatically well-structured
529 written text while the recalls comprise transcripts of casual unrehearsed speech.

530 We corrected for these representational artifacts by mean-centering the episode and recall trajec-
531 tories before performing analyses that involved measuring similarities between embedding vectors
532 within or across trajectories. Conceptually, mean-centering a given set of embedding vectors (i.e.,
533 subtracting the across-vector average from each individual vector) serves to remove representa-
534 tions that are common across all vectors under consideration, and instead highlight representations
535 that *vary* between them. For example, consider the sequence of embedding vectors comprising
536 the trajectory of the *Atlanta* episode. Common across vectors within this sequence will be latent
537 signals driven by how we chose to format the annotation text passed to the embedding model, the
538 annotator’s particular idiolect (i.e., idiosyncratic vocabulary and writing style), and any high-level
539 overarching themes of the episode. *Not* common across these vectors are the finer-grained semantic
540 representations that distinguish different events, scenes, and actions throughout the episode from
541 each other. We found that removing these non-discriminative signals (to the extent they behave as
542 additive offsets) via mean-centering helped reduce the embeddings’ anisotropy (i.e., increased the
543 range of measured between-vector similarities) and improved our ability to segment the episode
544 trajectory into meaningful events. Similarly, mean-centering both the episode and recalls before
545 relating their embedding representations helped mitigate broad-scale representational differences
546 driven by their diverging registers, and enabled us to better relate them based on event-specific,
547 content-driven patterns.

548 Our approach to mean-centering the recalls was guided by the particular latent representation
549 we wanted to remove versus retain. In segmenting each participant’s recall trajectory from each
550 session into events (see *Event-segmentation*), our goal was to discriminate between embedding
551 vectors strictly within a single recall sequence in order to identify temporally transient content-
552 related patterns. We therefore centered each recall trajectory (in addition to the episode) on its own
553 centroid to remove non-discriminative representations (e.g., register, broad-scale thematic content)
554 and maximize the contrast between vectors within that single sequence. However, the analyses
555 shown in Figures **REF TO DTW FIG** and **REF TO RDP FIG** entailed comparing sequences of

recall events not only to the episode, but also to each other, both within- and across-participants, to examine whether within-participant similarities tended to exceed across-participant similarities. Here our goal was to remove representational artifacts that could work to bias our results towards significance. Specifically, participants’ idiolects (i.e., particular speaking styles and sets of words with which they describe the episode) will influence the embedding representations of their recalls in ways that tend to be similar between their immediate and delayed recalls, but different across participants. To correct for this confound, we jointly centered each participant’s set of immediate and delayed recall events on the two sequences’ shared centroid (i.e., the average vector across all events from both sessions). This served to reduce the recall embeddings’ anisotropy (by removing signals driven by high-level themes of the episode they recalled) and mitigate the effects of the spoken-recall register and participant-specific idiolect (both of which are shared between a participant’s recalls from the two sessions) while preserving any genuine drift in the high-level content of each participant’s recall over the delay (which per-session mean-centering would instead discard).

Event-segmentation

After constructing trajectories for the episode and each participant’s immediate and delayed recall, we used hidden Markov models (HMMs; Baum and Petrie, 1966; Rabiner, 1989) to segment each trajectory into a sequence of discrete events. Given a number of states K , an HMM models an observed time series as a progression through K hidden states, each of which has a characteristic signature in the data (i.e., a pattern that timepoints belonging to that state tend to exhibit) and tends to persist across timepoints with occasional transitions rather than fluctuating freely. Because states persist, inferring the most probable state at each timepoint segments the observed time series into a sequence of events—i.e., spans of contiguous timepoints with the same most-probable state.

We segmented the episode trajectory into events using the HMM implemented in the BrainIAK toolbox (Kumar et al., 2021), which jointly learns a mean pattern for each of K states and the set of transitions between them via the Baum-Welch algorithm (Baum et al., 1970), with the added constraint that each state is visited exactly once in a strict left-to-right order (Baldassano et al., 2017).

583 To select an appropriate number of states (i.e., events) for the episode, we used an optimization
584 procedure we developed in prior work (Heusser et al., 2021) that aims to identify the single
585 timescale of event representations emphasized most heavily in the episode’s narrative structure.
586 We began by segmenting the (mean-centered) episode trajectory into K events for each K between
587 2 and 50, inclusive. For each candidate segmentation, we then computed the cosine similarity
588 between each pair of timepoints assigned to the same event, and between each pair of timepoints
589 assigned to different events. When constructing these distributions of within- and across-event
590 similarities, we restricted comparisons to pairs of timepoints separated by no more than the
591 duration of the longest event in the candidate segmentation, so that the within- and across-
592 event distributions were computed over comparable temporal distances and the across-event
593 distribution was not dominated by distant timepoint pairs that we would expect to be dissimilar
594 regardless of event structure. We then selected the candidate segmentation that maximized the
595 Wasserstein distance between these within- and across-event similarity distributions. We found
596 that the maximally distinctive and internally self-similar segmentation was yielded by $K = 20$
597 events. **(TODO: ADD SUPP FIG SHOWING WASSERSTEIN DIST AS FUNCTION OF K)**
598 Finally, we averaged the embedding vectors over timepoints between each pair of adjacent event
599 boundaries to obtain a sequence of stable event-average representations for the episode.

600 To segment participants’ recall trajectories into sequences of events, we developed a novel
601 HMM-driven procedure that grounded the recalls directly against the content of the episode
602 participants viewed and recounted. Essentially, rather than re-learning separate sets of state
603 representations from each participant’s recalls (which tend to be shorter, noisier approximations
604 of the episode content), we leveraged the state representations already learned from the episode to
605 identify transitions between the episode events a given participant was (most probably) describing
606 at a given point in time. First, we mean-centered the sequence of event-average embedding
607 vectors for the episode, and each participant’s sequence of per-window embedding vectors from
608 each session, separately. For each given recall trajectory, we then computed the cosine similarity
609 between the embedding vector for each recall window and the event-average embedding vector
610 for each of the 20 episode events. This yielded a number-of-windows by number-of-episode-

611 events (20) similarity matrix. We then treated the episode events as the hidden states of a hidden
612 Markov model with fixed parameters, and treated each recall window’s profile of similarities
613 to the episode events, z-scored across all window-event pairs, as its distribution of emission
614 scores over states. We imposed an ergodic transition structure (i.e., we allowed transitions from
615 any state to any other state, since participants may recall the episode events out of order, and
616 allowed states to be re-entered multiple times, since participants may revisit earlier events in
617 their recalls) governed by separate additive transition penalties for forward-one transitions and
618 all other transitions. Specifically, we imposed no penalty for remaining in the current state, a
619 small penalty (1.5 normalized emissions units) for transitioning from a given state to the next
620 (i.e., transitioning from recalling episode event n to recalling episode event $n + 1$), and a larger
621 penalty (3.0 normalized emissions units) for all other transitions. Finally, we decoded the most
622 probable sequence of episode-event labels using the Viterbi algorithm (Viterbi, 1967), collapsed
623 runs of consecutive identically labeled recall windows into events, and averaged the embedding
624 vectors over windows within each event. This procedure yielded, for each participant’s recall from
625 each session, a sequence of event-average embedding vectors and a set of labels denoting which
626 episode event each recall event corresponded to.

627 **Comparing trajectory structures**

628 We compared the geometric structures of the episode and recall event-sequences using dynamic
629 time warping (DTW; Sakoe and Chiba, 1978; Vintsyuk, 1968). The DTW algorithm finds the optimal
630 order-preserving alignment between two sequences of vectors with potentially different lengths or
631 sampling rates by locally “stretching” and “compressing” them in time to maximize their overall
632 similarity. This alignment yields a “warp path” that maps each vector in each sequence onto
633 one or more consecutive vectors in the other sequence with the minimum total distance between
634 matched vectors. The summed or average distance between matched vectors along the warp path
635 (the “warp cost”) is commonly taken as a measure of the two sequences’ structural dissimilarity.
636 We computed the warp cost as the average cosine distance between DTW-matched events so that
637 the cost values would be comparable across comparisons involving event-sequences of different

lengths (whereas summed distance necessarily grows with sequence length).

Data availability

All of the data analyzed in this manuscript may be found at <https://github.com/ContextLab/memory-dynamics>.

Code availability

All of the code for running our experiment and carrying out the analyses may be found at <https://github.com/ContextLab/memory-dynamics>.

References

- Anderson, A. J., McDermott, K., Rooks, B., Heffner, K. L., Dodell-Feder, D., and Lin, F. V. (2020). Decoding individual identity from brain activity elicited in imagining common experiences. *Nature Communications*, 11(1):5916.
- Anderson, R. C. and Pichert, J. W. (1978). Recall of previously unrecalable information following a shift in perspective. *Journal of Verbal Learning and Verbal Behavior*, 17(1):1–12.
- Antony, J., Lozano, A., Dhoat, P., Chen, J., and Bennion, K. (2024). Causal and chronological relationships predict memory organization for nonlinear narratives. *Journal of Cognitive Neuroscience*, 36(11):2368–2385.
- Baldassano, C., Chen, J., Zadbood, A., Pillow, J. W., Hasson, U., and Norman, K. A. (2017). Discovering event structure in continuous narrative perception and memory. *Neuron*, 95(3):709–721.
- Bartlett, F. C. (1932). *Remembering: A study in experimental and social psychology*. Oxford University Press.

- 659 Baum, L. E. and Petrie, T. (1966). Statistical inference for probabilistic functions of finite state
660 Markov chains. *Annals of Mathematical Statistics*, 37(6):1554–1563.
- 661 Baum, L. E., Petrie, T., Soules, G., and Weiss, N. (1970). A maximization technique occurring in the
662 statistical analysis of probabilistic functions of Markov chains. *Annals of Mathematical Statistics*,
663 41(1):164–171.
- 664 Bergman, E. T. and Roediger, H. L. (1999). Can Bartlett’s repeated reproduction experiments be
665 replicated? *Memory and Cognition*, 27(6):937–947.
- 666 Blei, D. M., Ng, A. Y., and Jordan, M. I. (2003). Latent dirichlet allocation. *Journal of Machine*
667 *Learning Research*, 3:993–1022.
- 668 Bridge, D. J. and Paller, K. A. (2012). Neural correlates of reactivation and retrieval-induced
669 distortion. *The Journal of Neuroscience*, 32(35):12144–12151.
- 670 Chan, J. C. K., Thomas, A. K., and Bulevich, J. B. (2009). Recalling a witnessed event increases
671 eyewitness suggestibility. *Psychological Science*, 20(1):66–73.
- 672 Chen, A., Zayats, V., Walker, D., and Padfield, D. (2022). Teaching BERT to wait: balancing accuracy
673 and latency for streaming disfluency detection. In Carpuat, M., de Marneffe, M. C., and Ruiz,
674 I. V. M., editors, *Proceedings of the 2022 Conference of the North American Chapter of the Association*
675 *for Computational Linguistics: Human Language Technologies*, pages 827–838, Seattle, United States.
676 Association for Computational Linguistics.
- 677 Chen, J., Leong, Y. C., Honey, C. J., Yong, C. H., Norman, K. A., and Hasson, U. (2017). Shared
678 memories reveal shared structure in neural activity across individuals. *Nature Neuroscience*,
679 20(1):115.
- 680 de Leeuw, J. R. (2015). jsPsych: A JavaScript library for creating behavioral experiments in a web
681 browser. *Behavior Research Methods*, 47(1):1–12.
- 682 Diamond, N. B. and Levine, B. (2020). Linking detail to temporal structure in naturalistic-event
683 recall. *Psychological Science*, 31(12):1557–1572.

- 684 Fenerci, C., Cheng, Z., Addis, D. R., Bellana, B., and Sheldon, S. (2025). Studying memory narratives
685 with natural language processing. *Trends in Cognitive Sciences*, 29(6):516–525.
- 686 Finn, E. S., Glerean, E., Khojandi, A. Y., Nielson, D., Molfese, P. J., Handwerker, D. A., and
687 Bandettini, P. A. (2020). Idiosynchrony: from shared responses to individual differences during
688 naturalistic neuroimaging. *NeuroImage*, 215:116828.
- 689 Finn, E. S., Shen, X., Scheinost, D., Rosenberg, M. D., Huang, J., Chun, M. M., Papademetris, X.,
690 and Constable, R. T. (2015). Functional connectome fingerprinting: identifying individuals using
691 patterns of brain connectivity. *Nature Neuroscience*, 18:1664–1671.
- 692 Fitzpatrick, P. C., Heusser, A. C., and Manning, J. R. (2026). Text embedding models yield detailed
693 conceptual knowledge maps derived from short multiple-choice quizzes. *Nature Communications*,
694 17(1):2955.
- 695 Gilboa, A. and Moscovitch, M. (2021). No consolidation without representation: correspon-
696 dence between neural and psychological representations in recent and remote memory. *Neuron*,
697 109(14):2239–2255.
- 698 Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., Halpern, D.,
699 Hamrick, J. B., and Chan, P. (2016). psiTurk: an open-source framework for conducting replicable
700 behavioral experiments online. *Behavior Research Methods*, 48(3):829–842.
- 701 Heusser, A. C., Fitzpatrick, P. C., and Manning, J. R. (2021). Geometric models reveal behavioural
702 and neural signatures of transforming experiences into memories. *Nature Human Behaviour*,
703 5(7):905–919.
- 704 Hopfield, J. J. (1982). Neural networks and physical systems with emergent collective computa-
705 tional abilities. *Proceedings of the National Academy of Sciences, USA*, 79:2554–2558.
- 706 Hupbach, A., Gomez, R., Hardt, O., and Nadel, L. (2007). Reconsolidation of episodic memories:
707 a subtle reminder triggers integration of new information. *Learning and Memory*, 14:47–53.

- 708 Hubbach, A., Gomez, R., and Nadel, L. (2009). Episodic memory reconsolidation: updating or
709 source confusion? *Memory*, 17(5):502–510.
- 710 Karpicke, J. D. and Roediger, H. L. (2007). Repeated retrieval during learning is the key to long-term
711 retention. *Journal of Memory and Language*, 57:151–162.
- 712 Kipp, M. (2001). Anvil-a generic annotation tool for multimodal dialogue. In *Proceedings of the*
713 *Conference on Speech Communication and Technology*.
- 714 Kumar, M., Anderson, M. J., Antony, J. W., Baldassano, C., Brooks, P. P., Cai, M. B., Chen, P.-H.,
715 Ellis, C. T., Henselman-Petrusek, G., Huberdeau, D., Hutchinson, J. B., Li, Y. P., Lu, Q., Manning,
716 J. R., Mennen, A. C., Nastase, S. A., Richard, H., Schapiro, A. C., Schuck, N. W., Shvartsman, M.,
717 Sundaram, N., Suo, D., Turek, J. S., Turner, D., Vo, V. A., Wallace, G., Wang, Y., Williams, J. A.,
718 Zhang, H., Zhu, X., Capota, M., Cohen, J. D., Hasson, U., Li, K., Ramadge, P. J., Turk-Browne,
719 N. B., Willke, T. L., and Norman, K. A. (2021). BrainIAK: the brain image analysis kit. *Aperture*,
720 pages doi.org/10.52294/31bb5b68–2184–411b–8c00–a1dacb61e1da.
- 721 Lee, H. and Chen, J. (2022). Predicting memory from the network structure of naturalistic events.
722 *Nature Communications*, 13(1):4235.
- 723 Lifanov, J., Linde-Domingo, J., and Wimber, M. (2021). Feature-specific reaction times reveal a
724 semanticisation of memories over time and with repeated remembering. *Nature Communications*,
725 12(1):3177.
- 726 Manning, J. R. (2021). Episodic memory: mental time travel or a quantum “memory wave”
727 function? *Psychological Review*, 128(4):711–725.
- 728 Manning, J. R. (2024). Context reinstatement. In Kahana, M. J. and Wagner, A. D., editors, *The*
729 *Oxford Handbook of Human Memory*. Oxford University Press.
- 730 Manning, J. R., Notaro, G. M., Chen, E., and Fitzpatrick, P. C. (2022). Fitness tracking reveals task-
731 specific associations between memory, mental health, and physical activity. *Scientific Reports*,
732 12(1):13822.

- 733 Marsh, E. J. (2007). Retelling is not the same as recalling. *Current Directions in Psychological Science*,
734 16(1):16–20.
- 735 Melega, G., Samson, K., Lee, H., and Renoult, L. (2025). The semantic structure of events consis-
736 tently influences episodic memory recall over time in young and older adults. *Scientific Reports*,
737 15(1):41987.
- 738 Nader, K. and Hardt, O. (2009). A single standard for memory: the case for reconsolidation. *Nature*
739 *Reviews Neuroscience*, 10(3):224–234.
- 740 Nader, K., Schafe, G. E., and LeDoux, J. E. (2000). Fear memories require protein synthesis in the
741 amygdala for reconsolidation after retrieval. *Nature*, 406(6797):722–726.
- 742 Nguyen, M., Vanderwal, T., and Hasson, U. (2019). Shared understanding of narratives is correlated
743 with shared neural response. *NeuroImage*, 184:161–170.
- 744 Oedekoven, C. S. H., Keidel, J. L., Berens, S. C., and Bird, C. M. (2017). Reinstatement of memory
745 representations for lifelike events over the course of a week. *Scientific Reports*, 7(1):14305.
- 746 Palombo, D. J., Sheldon, S., and Levine, B. (2018). Individual differences in autobiographical
747 memory. *Trends in Cognitive Sciences*, 22(7):583–597.
- 748 Qwen Team (2026). Qwen3.5: towards native multimodal agents. [https://qwen.ai/blog?id=](https://qwen.ai/blog?id=qwen3.5)
749 [qwen3.5](https://qwen.ai/blog?id=qwen3.5).
- 750 Rabiner, L. R. (1989). A tutorial on Hidden Markov Models and selected applications in speech
751 recognition. *Proceedings of the IEEE*, 77(2):257–286.
- 752 Raccah, O., Chen, P., Gureckis, T. M., Poeppel, D., and Vo, V. A. (2024). The “Naturalistic Free
753 Recall” dataset: four stories, hundreds of participants, and high-fidelity transcriptions. *Scientific*
754 *Data*, 11(1):1317.
- 755 Radvansky, G. A. and Zacks, J. M. (2017). Event boundaries in memory and cognition. *Current*
756 *Opinion in Behavioral Sciences*, 17:133–140.

- 757 Roediger, H. L. and Karpicke, J. D. (2006). Test-enhanced learning. *Psychological Science*, 17(3):249–
758 255.
- 759 Sakoe, H. and Chiba, S. (1978). Dynamic programming algorithm optimization for spoken word
760 recognition. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 26(1):43–49.
- 761 Sava-Segal, C., Richards, C., Leung, M., and Finn, E. S. (2023). Individual differences in neural
762 event segmentation of continuous experiences. *Cerebral Cortex*, 33(13):8164–8178.
- 763 Schacter, D. L. (2012). Constructive memory: past and future. *Dialogues in Clinical Neurosciences*,
764 1:7–18.
- 765 Schacter, D. L., Guerin, S. A., and St Jacques, P. L. (2011). Memory distortion: an adaptive
766 perspective. *Trends in Cognitive Sciences*, 15(10):467–474.
- 767 Sederberg, P. B., Gershman, S. J., Polyn, S. M., and Norman, K. A. (2011). Human memory
768 reconsolidation can be explained using the temporal context model. *Psychonomic Bulletin and*
769 *Review*, 18(3):455–468.
- 770 Sekeres, M. J., Bonasia, K., St-Laurent, M., Pishdadian, S., Winocur, G., Grady, C., and Moscovitch,
771 M. (2016). Recovering and preventing loss of detailed memory: differential rates of forgetting
772 for detail types in episodic memory. *Learning and Memory*, 23(2):72–82.
- 773 Sekeres, M. J., Winocur, G., and Moscovitch, M. (2018). The hippocampus and related neocortical
774 structures in memory transformation. *Neuroscience Letters*, 680:39–53.
- 775 St Jacques, P. L. and Schacter, D. L. (2013). Modifying memory: selectively enhancing and updating
776 personal memories for a museum tour by reactivating them. *Psychological Science*, 24(4):537–543.
- 777 Stone, S., Crossett, J., Luker, T., Leligdon, L., Cowen, W., and Darabos, C. (2025). Dartmouth Chat
778 - deploying an open-source LLM stack at scale. In *Practice and Experience in Advanced Research*
779 *Computing 2025: the Power of Collaboration*, pages 1–5, New York, NY, USA. Association for
780 Computing Machinery.

781 Su, H., Zhang, M., Knight, C., Zhao, M., and Lee, H. (2026). Retention and transformation of internal
782 experiences in autobiographical memory narratives. *Communications Psychology*, 4(1):56.

783 Thorndike, E. L. (1913). *Educational psychology: the psychology of learning*, volume 2. Teachers
784 College, Columbia University.

785 Tversky, B. and Marsh, E. J. (2000). Biased retelling of events yields biased memories. *Cognitive*
786 *Psychology*, 40:1–38.

787 Vera, H. S., Dua, S., Zhang, B., Salz, D., Mullins, R., Panyam, S. R., Smoot, S., Naim, I., Zou, J.,
788 Chen, F., Cer, D., Lisak, A., Choi, M., Gonzalez, L., Sanseviero, O., Cameron, G., Ballantyne, I.,
789 Black, K., Chen, K., Wang, W., Li, Z., Martins, G., Lee, J., Sherwood, M., Ji, J., Wu, R., Zheng,
790 J., Singh, J., Sharma, A., Sreepathihalli, D., Jain, A., Elarabawy, A., Co, A., Doumanoglou, A.,
791 Samari, B., Hora, B., Potetz, B., Kim, D., Alfonseca, E., Moiseev, F., Han, F., Gomez, F. P., Ábrego,
792 G. H., Zhang, H., Hui, H., Han, J., Gill, K., Chen, K., Chen, K., Shanbhogue, M., Boratko, M.,
793 Suganthan, P., Duddu, S. M. K., Mariserla, S., Ariafer, S., Zhang, S., Zhang, S., Baumgartner, S.,
794 Goenka, S., Qiu, S., Dabral, T., Walker, T., Rao, V., Khawaja, W., Zhou, W., Ren, X., Xia, Y., Chen,
795 Y., Chen, Y., Dong, Z., Ding, Z., Visin, F., Liu, G., Zhang, J., Kenealy, K., Casbon, M., Kumar,
796 R., Mesnard, T., Gleicher, Z., Brick, C., Lacombe, O., Roberts, A., Yin, Q., Sung, Y., Hoffmann,
797 R., Warkentin, T., Joulin, A., Duerig, T., and Seyedhosseini, M. (2025). EmbeddingGemma:
798 powerful and lightweight text representations. *arXiv*.

799 Vintsyuk, T. K. (1968). Speech discrimination by dynamic programming. *Cybernetics*, 4(1):52–57.

800 Viterbi, A. (1967). Error bounds for convolutional codes and an asymptotically optimum decoding
801 algorithm. *IEEE Transactions on Information Theory*, 13(2):260–269.

802 Wang, W., Tur, G., Zheng, J., and Ayan, N. F. (2010). Automatic disfluency removal for improving
803 spoken language translation. In *2010 IEEE International Conference on Acoustics, Speech and Signal*
804 *Processing*, pages 5214–5217, Dallas, TX, USA.

805 Xu, X., Zhu, Z., Zheng, X., and Manning, J. R. (2024). Temporal asymmetries in inferring unobserved
806 past and future events. *Nature Communications*, 15(1):8502.

807 Ziman, K., Heusser, A. C., Fitzpatrick, P. C., Field, C. E., and Manning, J. R. (2018). Is automatic
808 speech-to-text transcription ready for use in psychological experiments? *Behavior Research*
809 *Methods*, 50:2597–2605.

810 **Acknowledgements**

811 INDIVIDUAL ACKNOWLEDGEMENTS PLACEHOLDER. Our work was supported in part by
812 NSF CAREER Award Number 2145172 to J.R.M. The content is solely the responsibility of the
813 authors and does not necessarily represent the official views of our supporting organizations.
814 The funders had no role in study design, data collection and analysis, decision to publish, or
815 preparation of the manuscript.

816 **Author contributions**

817 **Competing interests**

818 The authors declare no competing interests.