

1

TITLE PLACEHOLDER

2

Paxton C. Fitzpatrick¹, Alishba Tahir^{1, 2}, Jennifer Xu¹, and Jeremy R. Manning^{1, *}

¹Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755, USA

²ALISHBA AFFILIATION PLACEHOLDER
ALISHBA AFFILIATION PLACEHOLDER

*Corresponding author: Jeremy.R.Manning@Dartmouth.edu

3

Abstract

4

ABSTRACT PLACEHOLDER

5

Introduction

6

Results

7

Discussion

8

Methods

9

Our experimental protocol was approved by the Committee for the Protection of Human Subjects

10

at Dartmouth College. We obtained written informed consent from all participants prior to the

11

start of both experiment sessions.

Participants

We enrolled a total of 70 Dartmouth undergraduate students in our study. Participants received optional course credit for enrolling and were pre-screened to ensure they had not previously seen any episode of either TV show viewed during the experiment. No statistical method was used to predetermine sample size. Prior to data processing and analysis, we excluded a total of 17 participants from our dataset: 4 encountered technical issues with the experiment, 2 self-reported difficulty with the task due to illness, 4 spoke for less than 10 minutes during the immediate recall phase, and 7 did not return for the second experiment session. The data we analyzed comprised the remaining 53 participants.

At the start of the experiment, we asked each participant to complete a demographic survey that included questions about their age, sex, race, and ethnicity. Participants' ages ranged from 18 to 22 years (mean: 18.91 years; standard deviation: 1.01 years). A total of 41 participants reported their sex as female and 12 participants reported their sex as male. Participants reported their races as White (34 participants), Asian (16 participants), Black or African American (3 participants), American Indian or Alaska Native (1 participant), and other (1 participant). (Note that some participants selected multiple racial categories.) A total of 52 participants reported their ethnicity as "Not Hispanic or Latino" and 1 participant reported their ethnicity as "Hispanic or Latino."

We also asked participants their native spoken language; whether they had any hearing, speech, or color vision impairments; and about any current medications or recent injuries that could impact attention or memory. All participants reported that their native spoken language was English, that they had no hearing or speech impairments, and that they had normal color vision. A total of 52 participants reported no attention- or memory-relevant medications or injuries. One participant reported sustaining a concussion more than 3 years prior to the experiment and was not excluded from our dataset.

Additionally, before each of the two experiment sessions, we asked participants about their sleep, coffee consumption, and level of alertness. Participants reported having slept between 5 and 11 hours the night prior to session 1 (mean: 7.37 hours; standard deviation: 1.45 hours) and

39 between 5 and 11 hours the night prior to session 2 (mean: 7.39 hours; standard deviation: 1.23
40 hours). They reported having consumed between 0 and 3 cups of coffee on the day of session
41 1 (mean: 0.43 cups; standard deviation: 0.72 cups) and between 0 and 3 cups of coffee on the
42 day of session 2 (mean: 0.48 cups; standard deviation: 0.68 cups). Participants reported their
43 current levels of alertness on a 5-point Likert scale. At the start of session 1, 1 participant reported
44 feeling "Very sluggish," 14 reported feeling "A little sluggish," 18 reported feeling "Neutral," 15
45 reported feeling "A little alert," and 5 reported feeling "Very alert." At the start of session 2, 3
46 participants reported feeling "Very sluggish," 9 reported feeling "A little sluggish," 23 reported
47 feeling "Neutral," 17 reported feeling "A little alert," and 3 reported feeling "Very alert."

48 @jeremy: do we want to report class years and majors? And/or do we want to report post-experiment
49 survey responses (how engaging did you find the episode, how easy/difficult was it to follow, how well do you
50 think you recalled it, how well do you think you learned characters' names)?

51 Experiment

52 The experiment took place over two sessions lasting approximately one hour each, separated by
53 6–8 days (mean: 6.76 days; standard deviation: 0.73 days). Prior to the start of the experiment,
54 participants were assigned to one of two groups (A or B), alternating based on enrollment order.
55 During session 1, all participants first viewed the pilot episode of the FX series *Atlanta* ("The Big
56 Bang"; duration: 24 min 15 s). Immediately after viewing, participants were given up to one hour
57 to recount what happened in the episode to the best of their abilities. Following Chen et al. (2017),
58 participants were instructed to describe the episode in as much detail as possible and to try to
59 speak for a minimum of 10 minutes. They were told to try to recount the episode's events in their
60 original order, but that they were free to go back and describe earlier events should they realize
61 they had missed something. After the immediate recall phase, participants were then given up to
62 5 minutes to predict what they thought would happen in the show's next episode.

63 Upon returning for session 2, all participants were again given up to one hour to recount the
64 *Atlanta* episode they had viewed the week prior. Following the delayed recall phase, participants in
65 Group A then viewed the second episode of *Atlanta* ("Streets on Lock"; duration: 21 min 57 s) while

66 participants in Group *B* instead viewed the pilot episode of the FOX series *Arrested Development*
67 (“Pilot”; duration: 20 min 37 s). Both groups were then given up to one hour to recount the
68 just-watched episode in the same manner and given the same instructions as in session 1. We
69 implemented the experiment using the psiTurk platform (Gureckis et al., 2016) and jsPsych library
70 (de Leeuw, 2015), with a custom jsPsych plugin we developed in prior work for recording verbal
71 recall audio (Ziman et al., 2018). Here we analyze participants’ immediate and delayed recalls of
72 the *Atlanta* episode viewed in session 1.

73 **Stimulus annotation and recall transcription**

74 To capture the evolving semantic content of the *Atlanta* episode, we created detailed text annota-
75 tions of the on-screen content during each individual camera shot throughout the episode (mean
76 annotation duration: 3.05 s; standard deviation: 1.80 s). Specifically, we used ANVIL (Kipp, 2001)
77 to record, for each camera shot, a set of semantically relevant features. These included a brief de-
78 scription of physical occurrences and characters’ explicit actions, a brief description of characters’
79 internal states (e.g., emotions, motivations) and other implicit content, the names of all characters
80 appearing on screen, a transcription of any dialogue, the name(s) of the character(s) speaking, the
81 shot’s setting, whether the shot took place indoors or outdoors, and whether or not there was
82 music playing during the shot. This feature set was adapted from a similar annotation corpus
83 created by Chen et al. (2017), which we used in developing our trajectory-analysis framework in
84 prior work (Heusser et al., 2021).

85 We then transcribed each participant’s immediate and delayed verbal recall of the *Atlanta*
86 episode. Our transcription process consisted of three steps: an initial manual transcription, a
87 two-pass prose-formatting conversion with a large language model, and a final manual cleanup
88 and correction step. First, we manually transcribed each recall audio file collected during the
89 experiment. These initial transcripts reflected participants’ spoken recalls verbatim, without any
90 punctuation or additional formatting, and including disfluencies, stuttered repetitions, and other
91 artifacts common to unrehearsed speech. However, since transformer models (such as we used
92 to obtain text embeddings of participants’ recalls) are trained almost exclusively on written text,

93 structural conventions like sentence breaks, punctuation, and capitalization provide an important
94 signal these models leverage to contextual other tokens (e.g., words) and construct a representation
95 of the text’s overall meaning. Meanwhile, speech-native phenomena such as stuttered repetitions
96 and filler words (“like”, “um”, etc.) that are comparatively rare in written text contribute out-of-
97 distribution noise that can be expected to degrade embedding fidelity, and prior work has shown
98 that automated removal of these phenomena improves performance on downstream tasks that
99 assume clean, written-style input (Chen et al., 2022; Wang et al., 2010). We therefore used a large
100 language model (Qwen3.5 122b; Qwen Team, 2026) to convert participants’ recall transcripts to a
101 structure and style more typical of written text. We used the `langchain_dartmouth` Python package
102 (Stone et al., 2025) to implement this conversion as two sequential passes over each transcript. In the
103 first pass, we prompted the model to split the transcript into sentences at natural breaking points,
104 add punctuation including quotes around reported speech, remove clear non-word disfluencies
105 (“um,” “uh,” “hmm,” etc.), collapse stuttered verbatim repetitions, and mark self-corrections and
106 restarts with em-dashes. In the second pass, we prompted the model to remove all instances of
107 “like” used as a filler word. We removed filler “like” instances in a separate pass because of their
108 high frequency of occurrence in the transcripts and the additional effort to required to distinguish
109 them from meaningful uses of “like.” The full texts of our prompts for both passes can be found in
110 *Supplementary methods*. Finally, we manually reviewed all LLM-formatted transcripts to identify
111 and correct any errors, ensure sentence breaks were appropriate and consistent across participants,
112 confirm that no content was improperly added to or removed from the transcripts, and resolve
113 any ambiguities based on the prosody of participants’ original recall audio.

114 **Analysis**

115 **Statistics**

116 Wilcoxon tests were one-sided because our alternative hypothesis was directional (delayed matches
117 own immediate *better than* the episode, others’ immediates, etc.).

118 Constructing episode and recall trajectories

119 We adapted a method we developed in prior work (Heusser et al., 2021) to transform the *Atlanta*
120 episode and participant’s recalls of it into trajectories through a common representation space
121 using a text embedding model. Briefly, text embedding models take in natural language text
122 of an arbitrary length and output a fixed-length vector representation that captures the text’s
123 semantic meaning. Intuitively, an output vector of length n can be viewed as a coordinate in an
124 n -dimensional text embedding space. While specific model architectures and training paradigms
125 vary, modern text embedding models are broadly trained to output similar coordinate vectors (in
126 terms of cosine similarity, dot product, inverse Euclidean distance, or other geometric measures) for
127 semantically similar input texts, even when those texts share few or no specific words. We obtained
128 text embeddings for the episode and recalls using EmbeddingGemma (Vera et al., 2025), a pre-
129 trained 308-million-parameter encoder-only transformer model that outputs coordinate vectors in
130 a 768-dimensional embedding space.

131 To construct the trajectory for the *Atlanta* episode, we first converted our tabular annotations
132 of the content present during each camera shot to a plain-text format. Separately for each shot,
133 we concatenated the annotation text describing explicit on-screen occurrences and characters’
134 implicit internal states. We then appended “Speech: ” followed by the shot’s dialogue, “Character
135 speaking: ” followed by the name of the speaker, and “Setting: ” followed by the shot’s setting. If
136 any annotated fields were empty for a given shot (e.g., if no character was speaking at that time), we
137 omitted them and their prefixes from the concatenated text. Next, we used EmbeddingGemma to
138 transform the text snippet for each camera shot into a text embedding vector, resulting in a number-
139 of-shots (478) by number-of-features (768) matrix. Since the camera shots varied in duration, we
140 then resampled this series of embedding vectors to a consistent temporal resolution. We assigned
141 each vector to the timepoint midway between the start and end of its corresponding camera shot
142 (in seconds), and used linear interpolation to estimate a vector for each second of the episode. This
143 yielded a number-of-seconds (1,455) by number-of-features (768) matrix—i.e., a 768-dimensional
144 *trajectory*.

145 We then created similar trajectories for each participant’s immediate and delayed recall of the
146 episode. We began by splitting each recall transcript into a list of sentences, with the constraint
147 that splits could not occur within reported speech (i.e., if a participant recounted a multi-sentence
148 quote by a character in the episode, we treated the full quote as a single sentence-unit). We
149 then parsed these sentences into overlapping sliding windows that spanned up to 5 consecutive
150 sentences each. These windows grew to and shrank from their full size at the beginning and end of
151 a participant’s recall sequence—in other words, the first window contained only the first sentence
152 of the participant’s recall, the second contained the first two sentences, and so on. This ensured
153 that all sentences from a participant’s recall transcript appeared an equal number of times (5) in
154 their sequence of sliding windows. Finally, we used EmbeddingGemma to transform each sliding
155 window into a text embedding vector. Repeated for all recall transcripts, this yielded a separate
156 trajectory characterizing each participant’s recall from each of the two experiment sessions. The
157 number of samples in participants’ recall trajectories (i.e., the number of sliding windows parsed
158 from their recall transcripts) ranged from 61 to 308 (mean: 136.68; standard deviation: 49.70).

159 **Relating episode and recall embedding vectors**

160 In prior work using our trajectory-analysis framework (e.g., Fitzpatrick et al., 2026; Heusser et al.,
161 2021; Manning, 2021; Manning et al., 2022; Racciah et al., 2024), embeddings of the stimulus and
162 participants’ responses were constructed using a topic model (Latent Dirichlet Allocation; Blei
163 et al., 2003) trained solely on the stimulus content. Here we elected to instead use a pre-trained
164 transformer-based text embedding model due to the richer and more nuanced textual information
165 these models are able to capture. (For example, these context-sensitive models can disambiguate
166 the texts “Earn looked at Alfred” and “Alfred looked at Earn” in their representations, whereas
167 bag-of-words models like LDA would represent the two identically.) However, these richer rep-
168 resentations also introduce certain geometric nuisances we needed to account for in order to
169 effectively relate the contents of participants’ recalls to that of the episode they viewed. First,
170 because modern transformer models are pre-trained on large and diverse text corpora, the repre-
171 sentational capacities of their embedding spaces span an enormous breadth of semantic content.

172 Consequently, in preliminary analyses we found that the embeddings of our annotations and par-
173 ticipants’ recalls (which were all “about” the events, themes, and characters of a single *Atlanta*
174 episode) were strongly anisotropic—i.e., they all occupied the same relatively narrow angular
175 region of the model’s overall embedding space. Additionally, because transformer models are sen-
176 sitive to the syntactic structure of text, their embedding representations carry latent information
177 about the text’s linguistic register in addition to its semantic content. Accordingly, we found there
178 was a consistent angular offset between embeddings of the episode annotations and those of par-
179 ticipants’ recalls, reflecting the fact that the annotations comprise grammatically well-structured
180 written text while the recalls comprise transcripts of casual unrehearsed speech.

181 We corrected for these representational artifacts by mean-centering the episode and recall trajec-
182 tories before performing analyses that involved measuring similarities between embedding vectors
183 within or across trajectories. Conceptually, mean-centering a given set of embedding vectors (i.e.,
184 subtracting the across-vector average from each individual vector) serves to remove representa-
185 tions that are common across all vectors under consideration, and instead highlight representations
186 that *vary* between them. For example, consider the sequence of embedding vectors comprising
187 the trajectory of the *Atlanta* episode. Common across vectors within this sequence will be latent
188 signals driven by how we chose to format the annotation text passed to the embedding model, the
189 annotator’s particular idiolect (i.e., idiosyncratic vocabulary and writing style), and any high-level
190 overarching themes of the episode. *Not* common across these vectors are the finer-grained semantic
191 representations that distinguish different events, scenes, and actions throughout the episode from
192 each other. We found that removing these non-discriminative signals (to the extent they behave as
193 additive offsets) via mean-centering helped reduce the embeddings’ anisotropy (i.e., increased the
194 range of measured between-vector similarities) and improved our ability to segment the episode
195 trajectory into meaningful events. Similarly, mean-centering both the episode and recalls before
196 relating their embedding representations helped mitigate broad-scale representational differences
197 driven by their diverging registers, and enabled us to better relate them based on event-specific,
198 content-driven patterns.

199 Our approach to mean-centering the recalls was guided by the particular latent representation

we wanted to remove versus retain. In segmenting each participant’s recall trajectory from each session into events (see *Event-segmentation*), our goal was to discriminate between embedding vectors strictly within a single recall sequence in order to identify temporally transient content-related patterns. We therefore centered each recall trajectory (in addition to the episode) on its own centroid to remove non-discriminative representations (e.g., register, broad-scale thematic content) and maximize the contrast between vectors within that single sequence. However, the analyses shown in Figures **REF TO DTW FIG** and **REF TO RDP FIG** entailed comparing sequences of recall events not only to the episode, but also to each other, both within- and across-participants, to examine whether within-participant similarities tended to exceed across-participant similarities. Here our goal was to remove representational artifacts that could work to bias our results towards significance. Specifically, participants’ idiolects (i.e., particular speaking styles and sets of words with which they describe the episode) will influence the embedding representations of their recalls in ways that tend to be similar between their immediate and delayed recalls, but different across participants. To correct for this confound, we jointly centered each participant’s set of immediate and delayed recall events on the two sequences’ shared centroid (i.e., the average vector across all events from both sessions). This served to reduce the recall embeddings’ anisotropy (by removing signals driven by high-level themes of the episode they recalled) and mitigate the effects of the spoken-recall register and participant-specific idiolect (both of which are shared between a participant’s recalls from the two sessions) while retaining any genuine drift in the high-level content of each participant’s recall over the delay (which per-session mean-centering would instead discard).

Event-segmentation

lkjlkj

223 **2D path embedding**

224 **Data availability**

225 All of the data analyzed in this manuscript may be found at [https://github.com/ContextLab/memory-](https://github.com/ContextLab/memory-dynamics)
226 [dynamics](https://github.com/ContextLab/memory-dynamics).

227 **Code availability**

228 All of the code for running our experiment and carrying out the analyses may be found at
229 <https://github.com/ContextLab/memory-dynamics>.

230 **References**

- 231 Blei, D. M., Ng, A. Y., and Jordan, M. I. (2003). Latent dirichlet allocation. *Journal of Machine*
232 *Learning Research*, 3:993–1022.
- 233 Chen, A., Zayats, V., Walker, D., and Padfield, D. (2022). Teaching BERT to wait: Balancing accuracy
234 and latency for streaming disfluency detection. In Carpuat, M., de Marneffe, M. C., and Ruiz,
235 I. V. M., editors, *Proceedings of the 2022 Conference of the North American Chapter of the Association*
236 *for Computational Linguistics: Human Language Technologies*, pages 827–838, Seattle, United States.
237 Association for Computational Linguistics.
- 238 Chen, J., Leong, Y. C., Honey, C. J., Yong, C. H., Norman, K. A., and Hasson, U. (2017). Shared
239 memories reveal shared structure in neural activity across individuals. *Nature Neuroscience*,
240 20(1):115.
- 241 de Leeuw, J. R. (2015). jsPsych: A JavaScript library for creating behavioral experiments in a web
242 browser. *Behavior Research Methods*, 47(1):1–12.
- 243 Fitzpatrick, P. C., Heusser, A. C., and Manning, J. R. (2026). Text embedding models yield detailed

conceptual knowledge maps derived from short multiple-choice quizzes. *Nature Communications*, 17(1):2955.

Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., Halpern, D., Hamrick, J. B., and Chan, P. (2016). psiTurk: an open-source framework for conducting replicable behavioral experiments online. *Behavior Research Methods*, 48(3):829–842.

Heusser, A. C., Fitzpatrick, P. C., and Manning, J. R. (2021). Geometric models reveal behavioral and neural signatures of transforming experiences into memories. *Nature Human Behaviour*, 5:905–919.

Kipp, M. (2001). Anvil - a generic annotation tool for multimodal dialogue. In *Proceedings of the 7th European Conference on Speech Communication and Technology*, pages 1367–1370.

Manning, J. R. (2021). Episodic memory: mental time travel or a quantum “memory wave” function? *Psychological Review*, 128(4):711–725.

Manning, J. R., Notaro, G. M., Chen, E., and Fitzpatrick, P. C. (2022). Fitness tracking reveals task-specific associations between memory, mental health, and physical activity. *Scientific Reports*, 12(1):13822.

Qwen Team (2026). Qwen3.5: Towards native multimodal agents. <https://qwen.ai/blog?id=qwen3.5>.

Racah, O., Chen, P., Gureckis, T. M., Poeppel, D., and Vo, V. A. (2024). The “Naturalistic Free Recall” dataset: four stories, hundreds of participants, and high-fidelity transcriptions. *Scientific Data*, 11(1):1317.

Stone, S., Crossett, J., Luker, T., Leligdon, L., Cowen, W., and Darabos, C. (2025). Dartmouth Chat - deploying an open-source llm stack at scale. In *Practice and Experience in Advanced Research Computing 2025: The Power of Collaboration*, number 43, pages 1–5, New York, NY, USA. Association for Computing Machinery.

268 Vera, H. S., Dua, S., Zhang, B., Salz, D., Mullins, R., Panyam, S. R., Smoot, S., Naim, I., Zou, J.,
 269 Chen, F., Cer, D., Lisak, A., Choi, M., Gonzalez, L., Sanseviero, O., Cameron, G., Ballantyne, I.,
 270 Black, K., Chen, K., Wang, W., Li, Z., Martins, G., Lee, J., Sherwood, M., Ji, J., Wu, R., Zheng,
 271 J., Singh, J., Sharma, A., Sreepathihalli, D., Jain, A., Elarabawy, A., Co, A., Doumanoglou, A.,
 272 Samari, B., Hora, B., Potetz, B., Kim, D., Alfonseca, E., Moiseev, F., Han, F., Gomez, F. P., Ábrego,
 273 G. H., Zhang, H., Hui, H., Han, J., Gill, K., Chen, K., Chen, K., Shanbhogue, M., Boratko, M.,
 274 Suganthan, P., Duddu, S. M. K., Mariserla, S., Ariafer, S., Zhang, S., Zhang, S., Baumgartner, S.,
 275 Goenka, S., Qiu, S., Dabral, T., Walker, T., Rao, V., Khawaja, W., Zhou, W., Ren, X., Xia, Y., Chen,
 276 Y., Chen, Y., Dong, Z., Ding, Z., Visin, F., Liu, G., Zhang, J., Kenealy, K., Casbon, M., Kumar,
 277 R., Mesnard, T., Gleicher, Z., Brick, C., Lacombe, O., Roberts, A., Yin, Q., Sung, Y., Hoffmann,
 278 R., Warkentin, T., Joulin, A., Duerig, T., and Seyedhosseini, M. (2025). EmbeddingGemma:
 279 Powerful and lightweight text representations. *arXiv*.

280 Wang, W., Tur, G., Zheng, J., and Ayan, N. F. (2010). Automatic disfluency removal for improving
 281 spoken language translation. In *2010 IEEE International Conference on Acoustics, Speech and Signal*
 282 *Processing*, pages 5214–5217, Dallas, TX, USA.

283 Ziman, K., Heusser, A. C., Fitzpatrick, P. C., Field, C. E., and Manning, J. R. (2018). Is automatic
 284 speech-to-text transcription ready for use in psychological experiments? *Behavior Research*
 285 *Methods*, 50:2597–2605.

286 Acknowledgements

287 INDIVIDUAL ACKNOWLEDGEMENTS PLACEHOLDER. Our work was supported in part by
 288 NSF CAREER Award Number 2145172 to J.R.M. The content is solely the responsibility of the
 289 authors and does not necessarily represent the official views of our supporting organizations.
 290 The funders had no role in study design, data collection and analysis, decision to publish, or
 291 preparation of the manuscript.

292 **Author contributions**

293 **Competing interests**

294 The authors declare no competing interests.