

1

TITLE PLACEHOLDER

2

Paxton C. Fitzpatrick¹, Alishba Tahir^{1, 2}, and Jeremy R. Manning^{1, *}

¹Department of Psychological and Brain Sciences
Dartmouth College, Hanover, NH 03755, USA

²ALISHBA AFFILIATION PLACEHOLDER
ALISHBA AFFILIATION PLACEHOLDER

*Corresponding author: Jeremy.R.Manning@Dartmouth.edu

3

Abstract

4

ABSTRACT PLACEHOLDER

5

Introduction

6

Results

7

Discussion

8

Methods

9

Our experimental protocol was approved by the Committee for the Protection of Human Subjects

10

at Dartmouth College. We obtained written informed consent from all participants prior to the

11

start of both experiment sessions.

Participants

We enrolled a total of 70 Dartmouth undergraduate students in our study. Participants received optional course credit for enrolling and were pre-screened to ensure they had not previously seen any episode of either TV show viewed during the experiment. No statistical method was used to predetermine sample size. Prior to data processing and analysis, we excluded a total of 17 participants from our dataset: 4 encountered technical issues with the experiment, 2 self-reported difficulty with the task due to illness, 4 spoke for less than 10 minutes during the immediate recall phase, and 7 did not return for the second experiment session. The data we analyzed comprised the remaining 53 participants.

At the start of the experiment, we asked each participant to complete a demographic survey that included questions about their age, sex, race, and ethnicity. Participants' ages ranged from 18 to 22 years (mean: 18.91 years; standard deviation: 1.01 years). A total of 41 participants reported their sex as female and 12 participants reported their sex as male. Participants reported their races as White (34 participants), Asian (16 participants), Black or African American (3 participants), American Indian or Alaska Native (1 participant), and other (1 participant). (Note that some participants selected multiple racial categories.) A total of 52 participants reported their ethnicity as "Not Hispanic or Latino" and 1 participant reported their ethnicity as "Hispanic or Latino."

We also asked participants their native spoken language; whether they had any hearing, speech, or color vision impairments; and about any current medications or recent injuries that could impact attention or memory. All participants reported that their native spoken language was English, that they had no hearing or speech impairments, and that they had normal color vision. A total of 52 participants reported no attention- or memory-relevant medications or injuries. One participant reported sustaining a concussion more than 3 years prior to the experiment and was not excluded from our dataset.

Additionally, before each of the two experiment sessions, we asked participants about their sleep, coffee consumption, and level of alertness. Participants reported having slept between 5 and 11 hours the night prior to session 1 (mean: 7.37 hours; standard deviation: 1.45 hours) and

39 between 5 and 11 hours the night prior to session 2 (mean: 7.39 hours; standard deviation: 1.23
40 hours). They reported having consumed between 0 and 3 cups of coffee on the day of session
41 1 (mean: 0.43 cups; standard deviation: 0.72 cups) and between 0 and 3 cups of coffee on the
42 day of session 2 (mean: 0.48 cups; standard deviation: 0.68 cups). Participants reported their
43 current levels of alertness on a 5-point Likert scale. At the start of session 1, 1 participant reported
44 feeling "Very sluggish," 14 reported feeling "A little sluggish," 18 reported feeling "Neutral," 15
45 reported feeling "A little alert," and 5 reported feeling "Very alert." At the start of session 2, 3
46 participants reported feeling "Very sluggish," 9 reported feeling "A little sluggish," 23 reported
47 feeling "Neutral," 17 reported feeling "A little alert," and 3 reported feeling "Very alert."

48 @jeremy: do we want to report class years and majors? And/or do we want to report post-experiment
49 survey responses (how engaging did you find the episode, how easy/difficult was it to follow, how well do you
50 think you recalled it, how well do you think you learned characters' names)?

51 Experiment

52 The experiment took place over two sessions lasting approximately one hour each, separated by
53 6–8 days (mean: 6.76 days; standard deviation: 0.73 days). Prior to the start of the experiment,
54 participants were assigned to one of two groups (A or B), alternating based on enrollment order.
55 During session 1, all participants first viewed the pilot episode of the FX series *Atlanta* ("The Big
56 Bang"; duration: 24 min 15 s). Immediately after viewing, participants were given up to one hour
57 to recount what happened in the episode to the best of their abilities. Following Chen et al. (2017),
58 participants were instructed to describe the episode in as much detail as possible and to try to
59 speak for a minimum of 10 minutes. They were told to try to recount the episode's events in their
60 original order, but that they were free to go back and describe earlier events should they realize
61 they had missed something. After the immediate recall phase, participants were then given up to
62 5 minutes to predict what they thought would happen in the show's next episode.

63 Upon returning for session 2, all participants were again given up to one hour to recount the
64 *Atlanta* episode they had viewed the week prior. Following the delayed recall phase, participants in
65 Group A then viewed the second episode of *Atlanta* ("Streets on Lock"; duration: 21 min 57 s) while

66 participants in Group *B* instead viewed the pilot episode of the FOX series *Arrested Development*
67 (“Pilot”; duration: 20 min 37 s). Both groups were then given up to one hour to recount the
68 just-watched episode in the same manner and given the same instructions as in session 1. We
69 implemented the experiment using the psiTurk platform (Gureckis et al., 2016) and jsPsych library
70 (de Leeuw, 2015), with a custom jsPsych plugin we developed in prior work for recording verbal
71 recall audio (Ziman et al., 2018). Here we analyze participants’ immediate and delayed recalls of
72 the *Atlanta* episode viewed in session 1.

73 **Stimulus annotation and recall transcription**

74 To capture the evolving semantic content of the *Atlanta* episode, we created detailed text annota-
75 tions of the on-screen content during each individual camera shot throughout the episode (mean
76 annotation duration: 3.05 s; standard deviation: 1.80 s). Specifically, we used ANVIL (Kipp, 2001)
77 to record, for each camera shot, a set of semantically relevant features. These included a brief de-
78 scription of physical occurrences and characters’ explicit actions, a brief description of characters’
79 internal states (e.g., emotions, motivations) and other implicit content, the names of all characters
80 appearing on screen, a transcription of any dialogue, the name(s) of the character(s) speaking, the
81 shot’s setting, whether the shot took place indoors or outdoors, and whether or not there was
82 music playing during the shot. This feature set was adapted from a similar annotation corpus
83 created by Chen et al. (2017), which we used in developing our trajectory-analysis framework in
84 prior work (Heusser et al., 2021).

85 We then transcribed each participant’s immediate and delayed verbal recall of the *Atlanta*
86 episode. Our transcription process consisted of three steps: an initial manual transcription, a
87 two-pass prose-formatting conversion with a large language model, and a final manual cleanup
88 and correction step. First, we manually transcribed each recall audio file collected during the
89 experiment. These initial transcripts reflected participants’ spoken recalls verbatim, without any
90 punctuation or additional formatting, and including disfluencies, stuttered repetitions, and other
91 artifacts common to unrehearsed speech. However, since transformer models (such as we used
92 to obtain text embeddings of participants’ recalls) are trained almost exclusively on written text,

93 structural conventions like sentence breaks, punctuation, and capitalization provide an important
94 signal these models leverage to contextual other tokens (e.g., words) and construct a representation
95 of the text’s overall meaning. Meanwhile, speech-native phenomena such as stuttered repetitions
96 and filler words (“like”, “um”, etc.) that are comparatively rare in written text contribute out-of-
97 distribution noise that can be expected to degrade embedding fidelity, and prior work has shown
98 that automated removal of these phenomena improves performance on downstream tasks that
99 assume clean, written-style input (Chen et al., 2022; Wang et al., 2010). We therefore used a large
100 language model (Qwen3.5 122b; Qwen Team, 2026) to convert participants’ recall transcripts to a
101 structure and style more typical of written text. We used the `langchain_dartmouth` Python package
102 (Stone et al., 2025) to implement this conversion as two sequential passes over each transcript. In the
103 first pass, we prompted the model to split the transcript into sentences at natural breaking points,
104 add punctuation including quotes around reported speech, remove clear non-word disfluencies
105 (“um,” “uh,” “hmm,” etc.), collapse stuttered verbatim repetitions, and mark self-corrections and
106 restarts with em-dashes. In the second pass, we prompted the model to remove all instances of
107 “like” used as a filler word. We removed filler “like” instances in a separate pass because of their
108 high frequency of occurrence in the transcripts and the additional effort to required to distinguish
109 them from meaningful uses of “like.” The full texts of our prompts for both passes can be found in
110 *Supplementary methods*. Finally, we manually reviewed all LLM-formatted transcripts to identify
111 and correct any errors, ensure sentence breaks were appropriate and consistent across participants,
112 confirm that no content was improperly added to or removed from the transcripts, and resolve
113 any ambiguities based on the prosody of participants’ original recall audio.

114 **Analysis**

115 **Statistics**

116 Wilcoxon tests were one-sided because our alternative hypothesis was directional (delayed matches
117 own immediate *better than* the episode, others’ immediates, etc.).

Constructing episode and recall trajectories

We adapted a method we developed in prior work (Heusser et al., 2021) to transform the *Atlanta* episode and participant’s recalls of it into trajectories through a common representation space using a text embedding model. Briefly, text embedding models take in natural language text of an arbitrary length and output a fixed-length vector representation that captures the text’s semantic meaning. Intuitively, an output vector of length n can be viewed as a coordinate in an n -dimensional text embedding space. While specific model architectures and training paradigms vary, modern text embedding models are broadly trained to output similar coordinate vectors (in terms of cosine similarity, dot product, inverse Euclidean distance, or other geometric measures) for semantically similar input texts, even when those texts share few or no specific words in common. We obtained text embeddings for the episode and recalls using EmbeddingGemma (Vera et al., 2025), a pre-trained 308-million-parameter encoder-only transformer model that outputs coordinate vectors in a 768-dimensional embedding space.

To construct the trajectory for the *Atlanta* episode, we first converted our tabular annotations of the content during each camera shot to a plain-text format. Separately for each shot, we concatenated the annotation text describing explicit on-screen occurrences and characters’ implicit internal states. We then appended “Speech: ” followed by the shot’s dialogue, “Character speaking: ” followed by the name of the speaker, and “Setting: ” followed by the shot’s setting. If any annotated fields were empty for a given shot (e.g., if no character was speaking at that time), we omitted them and their prefixes from the concatenated text. Next, we used EmbeddingGemma to transform the text snippet for each camera shot into a text embedding vector, resulting in a number-of-shots (478) by number-of-features (768) matrix. Since the camera shots varied in duration, we then resampled this series of embedding vectors to a consistent temporal resolution. We assigned each vector to the timepoint midway between the start and end of its corresponding camera shot (in seconds), and used linear interpolation to estimate a vector for each second of the episode. This yielded a number-of-seconds (1455) by number-of-features (768) matrix—i.e., a 768-dimensional *trajectory*.

We then created similar trajectories for each participant’s immediate and delayed recall of the

episode. We began by splitting each recall transcript into a list of sentences, with the constraint that splits could not occur within reported speech (i.e., if a participant recounted a multi-sentence quote by a character in the episode, we treated the full quote as a single sentence-unit). We then parsed these sentences into overlapping sliding windows that spanned up to 5 consecutive sentences each. These windows grew to and shrank from their full size at the beginning and end of a participant's recall sequence—in other words, the first window contained only the first sentence of the participant's recall, the second contained the first two sentences, and so on. This ensured that all sentences from a participant's recall appeared an equal number of times (5) in their sequence of sliding windows. Finally, we used EmbeddingGemma to transform each sliding window into a text embedding vector. Repeated for all recall transcripts, this yielded a separate trajectory characterizing each participant's recall from each of the two experiment sessions. The number of samples in participants' recall trajectories (i.e., the number of sliding windows parsed from their recall transcripts) ranged from 61 to 308 with a mean of 136.68 and standard deviation of 49.70.

Event-segmentation

2D path embedding

Creating knowledge and learning map visualizations

Data availability

All of the data analyzed in this manuscript may be found at <https://github.com/ContextLab/memory-dynamics>.

166 Code availability

167 All of the code for running our experiment and carrying out the analyses may be found at
168 <https://github.com/ContextLab/memory-dynamics>.

169 References

- 170 Chen, A., Zayats, V., Walker, D., and Padfield, D. (2022). Teaching BERT to wait: Balancing accuracy
171 and latency for streaming disfluency detection. In Carpuat, M., de Marneffe, M. C., and Ruiz,
172 I. V. M., editors, *Proceedings of the 2022 Conference of the North American Chapter of the Association
173 for Computational Linguistics: Human Language Technologies*, pages 827–838, Seattle, United States.
174 Association for Computational Linguistics.
- 175 Chen, J., Leong, Y. C., Honey, C. J., Yong, C. H., Norman, K. A., and Hasson, U. (2017). Shared
176 memories reveal shared structure in neural activity across individuals. *Nature Neuroscience*,
177 20(1):115.
- 178 de Leeuw, J. R. (2015). jsPsych: A JavaScript library for creating behavioral experiments in a web
179 browser. *Behavior Research Methods*, 47(1):1–12.
- 180 Gureckis, T. M., Martin, J., McDonnell, J., Rich, A. S., Markant, D., Coenen, A., Halpern, D.,
181 Hamrick, J. B., and Chan, P. (2016). psiTurk: an open-source framework for conducting replicable
182 behavioral experiments online. *Behavior Research Methods*, 48(3):829–842.
- 183 Heusser, A. C., Fitzpatrick, P. C., and Manning, J. R. (2021). Geometric models reveal behavioral
184 and neural signatures of transforming experiences into memories. *Nature Human Behaviour*,
185 5:905–919.
- 186 Kipp, M. (2001). Anvil - a generic annotation tool for multimodal dialogue. In *Proceedings of the
187 7th European Conference on Speech Communication and Technology*, pages 1367–1370.
- 188 Qwen Team (2026). Qwen3.5: Towards native multimodal agents. [https://qwen.ai/blog?id=](https://qwen.ai/blog?id=qwen3.5)
189 [qwen3.5](https://qwen.ai/blog?id=qwen3.5).

- 190 Stone, S., Crossett, J., Luker, T., Leligdon, L., Cowen, W., and Darabos, C. (2025). Dartmouth
191 Chat - deploying an open-source llm stack at scale. In *Practice and Experience in Advanced
192 Research Computing 2025: The Power of Collaboration*, number 43, pages 1–5, New York, NY, USA.
193 Association for Computing Machinery.
- 194 Vera, H. S., Dua, S., Zhang, B., Salz, D., Mullins, R., Panyam, S. R., Smoot, S., Naim, I., Zou, J.,
195 Chen, F., Cer, D., Lisak, A., Choi, M., Gonzalez, L., Sanseviero, O., Cameron, G., Ballantyne, I.,
196 Black, K., Chen, K., Wang, W., Li, Z., Martins, G., Lee, J., Sherwood, M., Ji, J., Wu, R., Zheng,
197 J., Singh, J., Sharma, A., Sreepathihalli, D., Jain, A., Elarabawy, A., Co, A., Doumanoglou, A.,
198 Samari, B., Hora, B., Potetz, B., Kim, D., Alfonseca, E., Moiseev, F., Han, F., Gomez, F. P., Ábrego,
199 G. H., Zhang, H., Hui, H., Han, J., Gill, K., Chen, K., Chen, K., Shanbhogue, M., Boratko, M.,
200 Suganthan, P., Duddu, S. M. K., Mariserla, S., Ariafer, S., Zhang, S., Zhang, S., Baumgartner, S.,
201 Goenka, S., Qiu, S., Dabral, T., Walker, T., Rao, V., Khawaja, W., Zhou, W., Ren, X., Xia, Y., Chen,
202 Y., Chen, Y., Dong, Z., Ding, Z., Visin, F., Liu, G., Zhang, J., Kenealy, K., Casbon, M., Kumar,
203 R., Mesnard, T., Gleicher, Z., Brick, C., Lacombe, O., Roberts, A., Yin, Q., Sung, Y., Hoffmann,
204 R., Warkentin, T., Joulin, A., Duerig, T., and Seyedhosseini, M. (2025). EmbeddingGemma:
205 Powerful and lightweight text representations. *arXiv*.
- 206 Wang, W., Tur, G., Zheng, J., and Ayan, N. F. (2010). Automatic disfluency removal for improving
207 spoken language translation. In *2010 IEEE International Conference on Acoustics, Speech and Signal
208 Processing*, pages 5214–5217, Dallas, TX, USA.
- 209 Ziman, K., Heusser, A. C., Fitzpatrick, P. C., Field, C. E., and Manning, J. R. (2018). Is automatic
210 speech-to-text transcription ready for use in psychological experiments? *Behavior Research
211 Methods*, 50:2597–2605.

212 Acknowledgements

213 INDIVIDUAL ACKNOWLEDGEMENTS PLACEHOLDER. Our work was supported in part by
214 NSF CAREER Award Number 2145172 to J.R.M. The content is solely the responsibility of the

215 authors and does not necessarily represent the official views of our supporting organizations.
216 The funders had no role in study design, data collection and analysis, decision to publish, or
217 preparation of the manuscript.

218 **Author contributions**

219 **Competing interests**

220 The authors declare no competing interests.